

A Survey on Spiking Neural Network for Drone Flight Control

スパイキングニューラルネットワークを用いた
ドローン飛行制御に関する研究調査

平田 涼馬(Ryoma Hirata)
Fukuda lab. M1

Outline

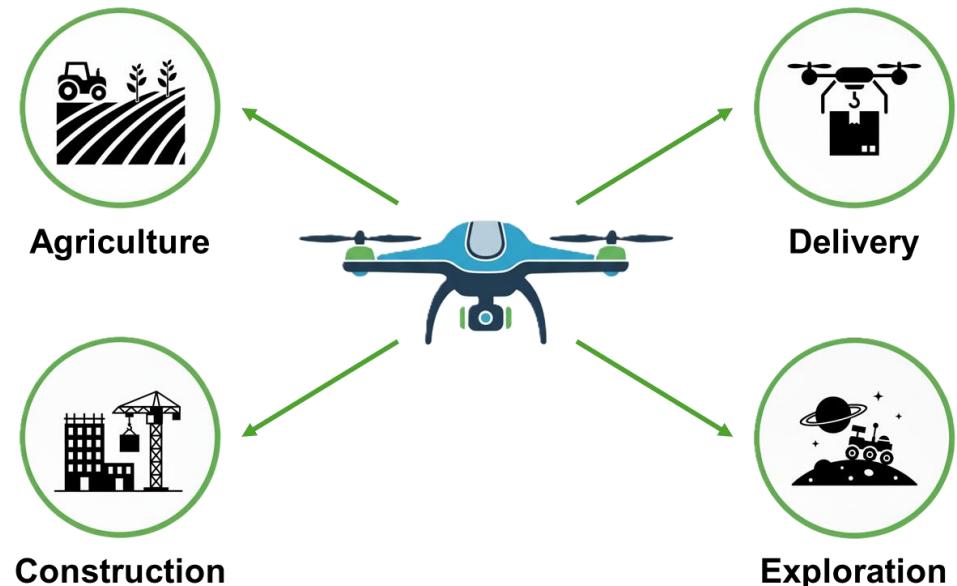
- 1. Introduction**
- 2. Neuromorphic Attitude Estimation and Control using IMU**
- 3. Neuromorphic Control using Event-based Vision**
- 4. Agile Navigation using Reinforcement Learning-Trained SNN**
- 5. Conclusion**

1. Introduction

- Benefits of Drones
 1. High mobility
 2. Low cost of introduction / operation

- Benefits of AI Control
 - **autonomous**
 - **performing complex tasks**

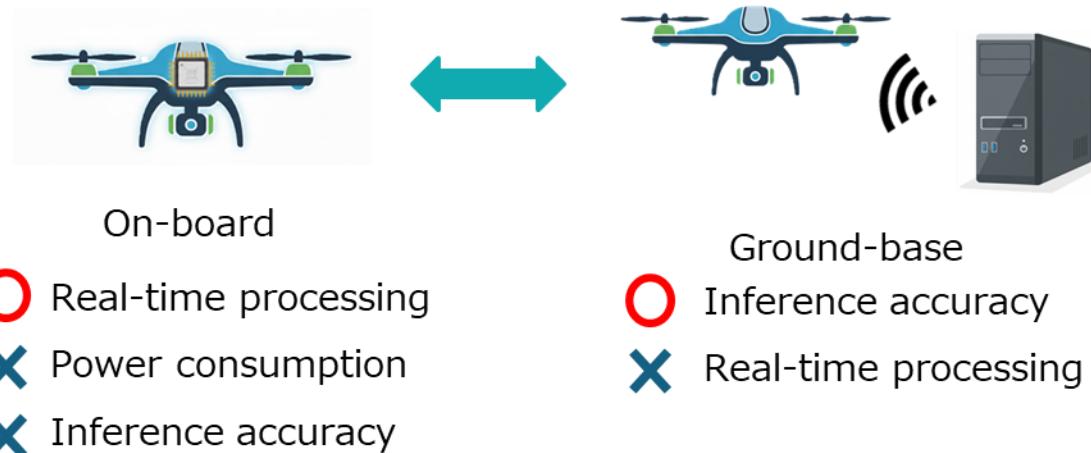
- Comparison: Traditional Control vs AI Control



features	Conventional Control (PID / Model base...)	AI Control
Adaptability	Low	High
Target Environment	Static	Dynamic, unpredictable, complex
Interpretability	High	Low
Reliability	High	Low

1. Introduction

Two methods of AI control for drones



Trade-off in On-board AI

Factor	Focusing on accuracy	Focusing on Flight Time
① Inference Accuracy	High	Low
② Model Size	Large	Small
③ Calculation amount	High	Low
④ Power Consumption	High	Low
⑤ Flight Time	Short	Long

- In the **Martian environment**, processing must be done **on-board**
 - There is too much latency in communication with Earth.
- NASA researchers propose using **Spiking Neural Networks** to control future Mars drones



Fig. Ingenuity^[1]

[1] Harbour, David AR, et al, "Martian Flight: Enabling Motion Estimation of NASA's Next-Generation Mars Flying Drone by Implementing a Neuromorphic Event-Camera and Explainable Fuzzy Spiking Neural Network Model." 2024 AIAA DASC/IEEE 43rd Digital Avionics Systems Conference (DASC). IEEE, 2024.

1. Introduction

Spiking Neural Network (SNN)

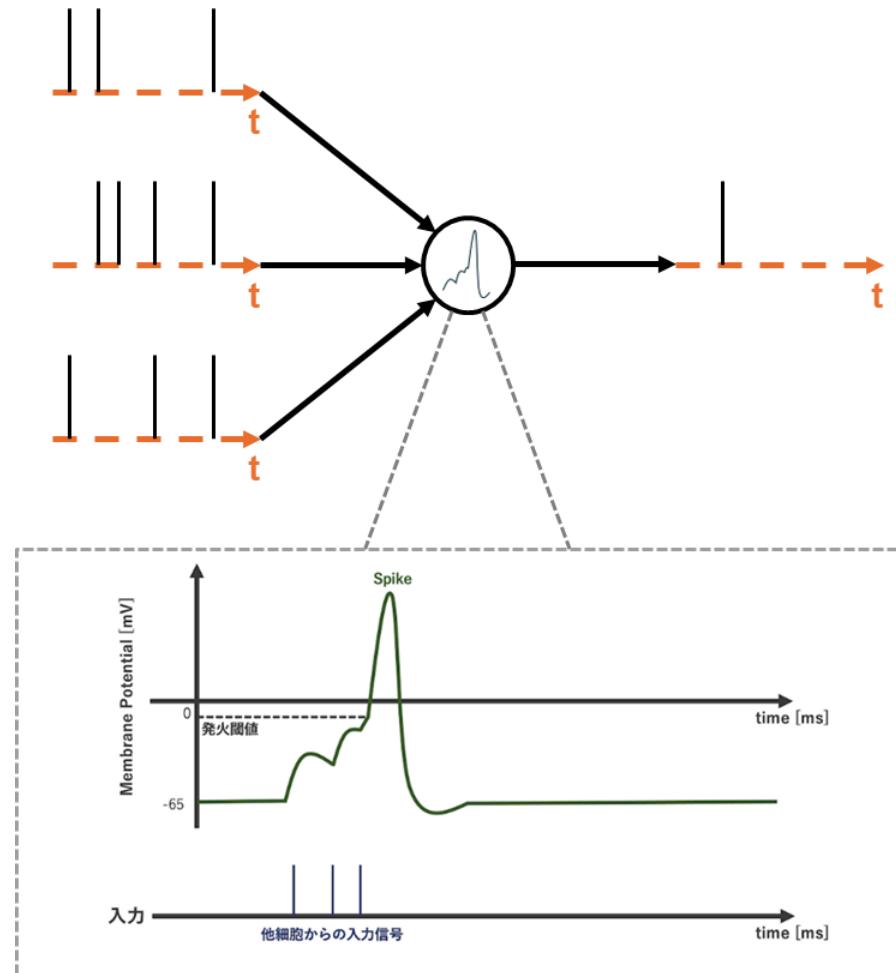


Fig. SNN Operation (Modified from [2])

LIFモデル

$$\tau_m \frac{dV(t)}{dt} = \underbrace{(-V(t) + E_{rest})}_{\text{Leak Term}} + \underbrace{I(t)}_{\text{Input Current}}$$

τ_m : Membrane Time Constant
 $I(t)$: Input Current
 $V(t)$: Membrane Potential
 E_{rest} : Resting Membrane Potential

- Low Power Consumption
- Low Latency Response
- ✗ Training Difficulty
- ✗ Hardware Dependency

[2] ゼロから学ぶスパイキングニューラルネットワーク <https://snn.hirlab.net/?s=2/>

1. Introduction

- This investigation focused on two applications:
low-level control (attitude control) and **path planning**

Application	Sensor	SNN output	Platform	Introduction paper
Low-level control (Attitude control)	IMU	Attitude / control	Microprocessor	Paper 1
Low-level control (Attitude control)	Event camera	Optical flow	Neuromorphic processor	Paper 2
Path planning	-	Attitude	Simulator	Paper 3

2. Neuromorphic Attitude Estimation and Control using IMU

Paper1

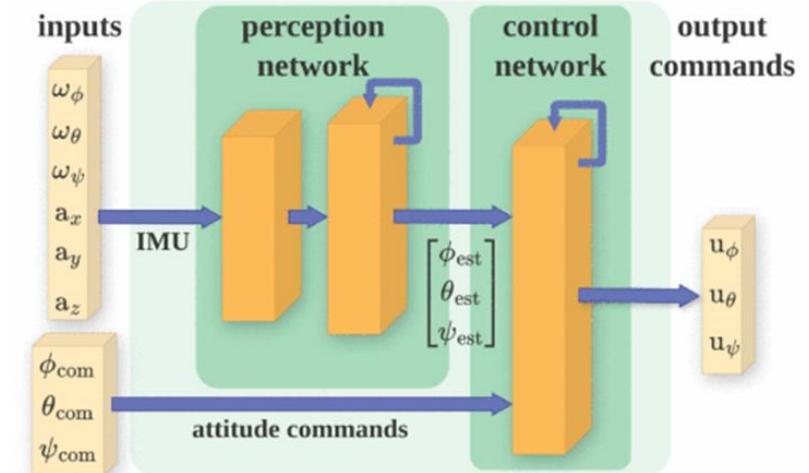
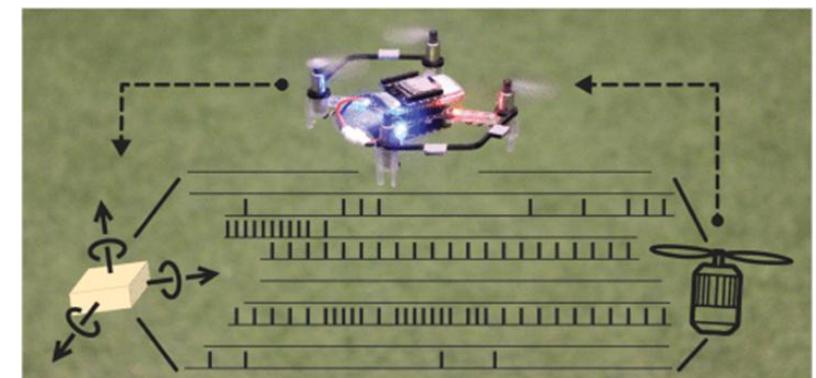
S. Stroobants, C. De. Wagter, and G. CHE. De. Croon,
 "Neuromorphic attitude estimation and control",
 IEEE Robotics and Automation Letters, (2025).

■ Key point

- This paper propose an **SNN model** that outputs **low-level control** commands from **IMU** inputs.

■ Network Architecture

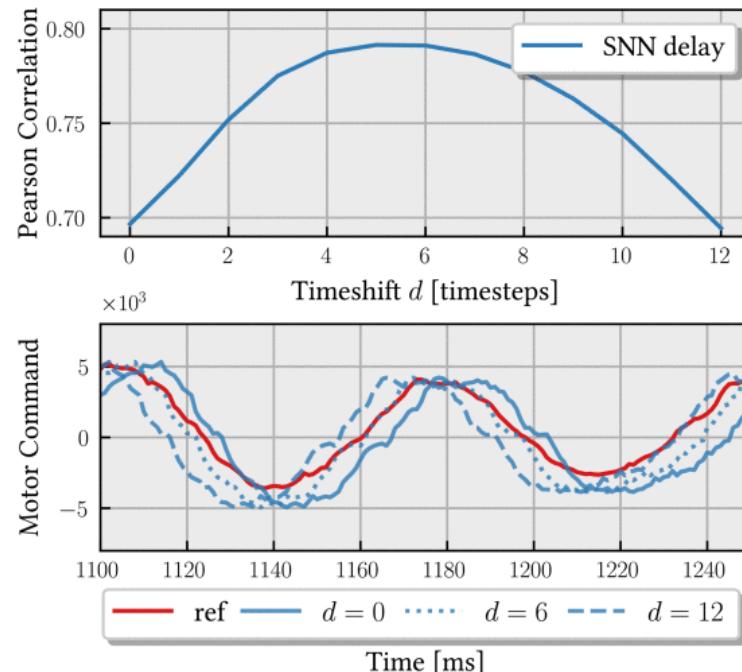
- Divided into two subnetworks
 1. Perception network (**SNN**)
 - Input : [Angular velocity, Acceleration]
 - Output : [Attitude (Euler kernel)]
 2. Control network (**SNN**)
 - Input : [Attitude (estimation + target)]
 - Output : [**Torque**]



2. Neuromorphic Attitude Estimation and Control using IMU

Training techniques

1. Handling SNN-specific delays



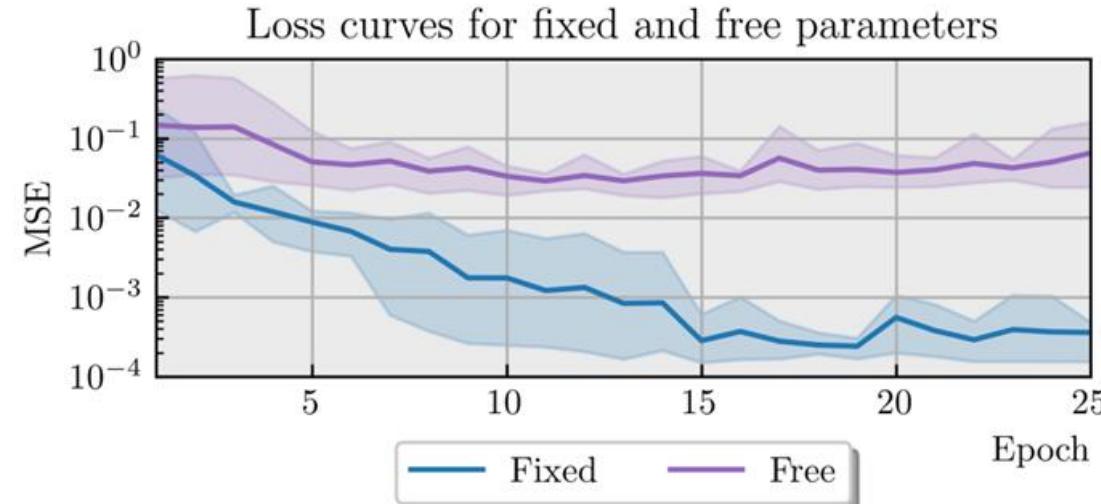
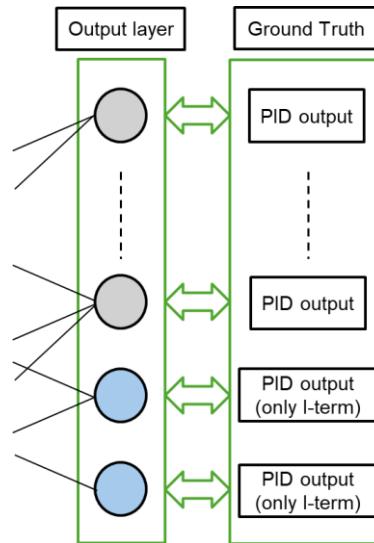
- Shift the ground truth labels by 6 steps

2. Data augmentation

- Flying using a SNN trained with the initial training data
- Gathering data by adding random disturbances to the PID controller

Training : Integral control compensation

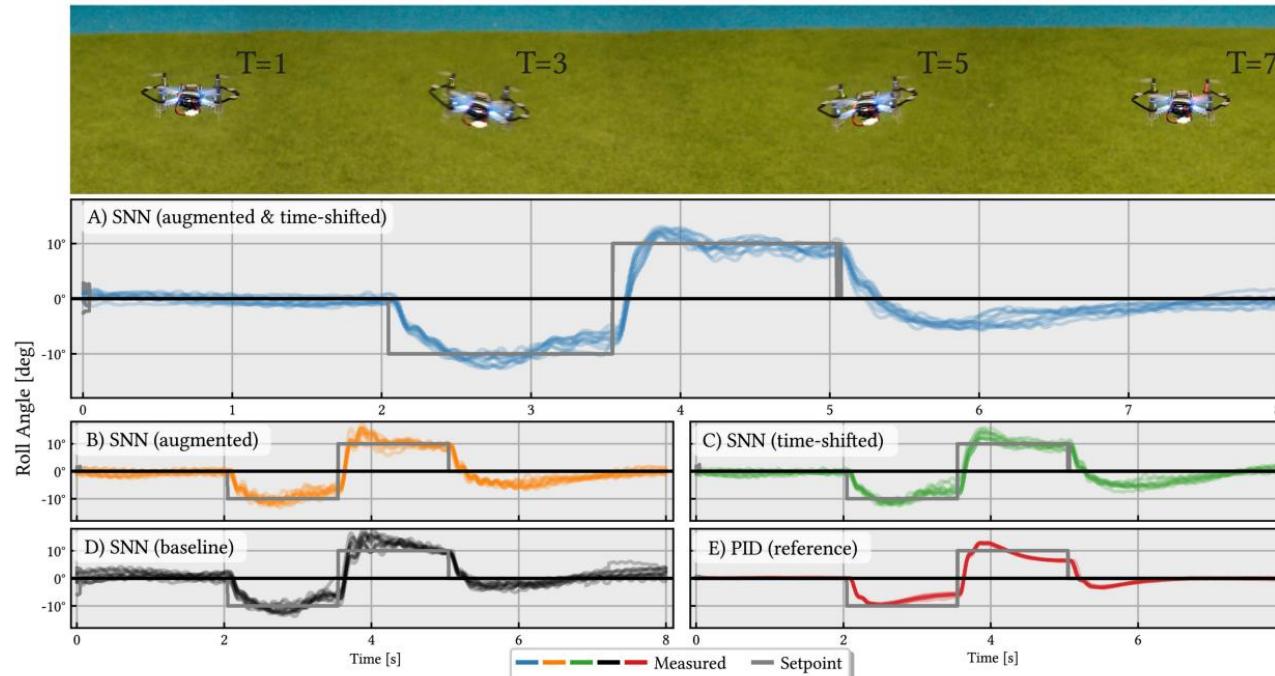
- Learning integral control of PID is difficult.
 - Learning PD control takes priority
 - membrane potential of neurons can't retain information for a long time due to **leakage**



- Fixing the leak parameters of 10 neurons to integrate input in the membrane potential
- Set the GT labels for those neurons to I-term of PID control

2. Neuromorphic Attitude Estimation and Control using IMU

Result : Roll angle control



Definition of RMSE
(Root Mean Square Error)

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=0}^{n-1} (y_i - \hat{y}_i)^2}$$

y_i : Ground truth,
 \hat{y}_i : Predicted value,
 n : total data count

Controller	RMSE	avg. SD	avg. RT
SNN (time-shifted & augm.)	3.03°	0.77	145ms
SNN (augmented)	3.10°	0.95	130ms
SNN (time-shifted)	3.24°	0.92	145ms
SNN (baseline)	3.14°	1.16	135ms
PID	2.67°	0.23	125ms

Note that the PID receives the estimated attitude as input, while the SNN needs to calculate this internally.

PID has higher tracking accuracy

- SNN's performance is **comparable** when considering that it is **also estimating attitude**

3. Neuromorphic Control using Event-based Vision

10

Paper2

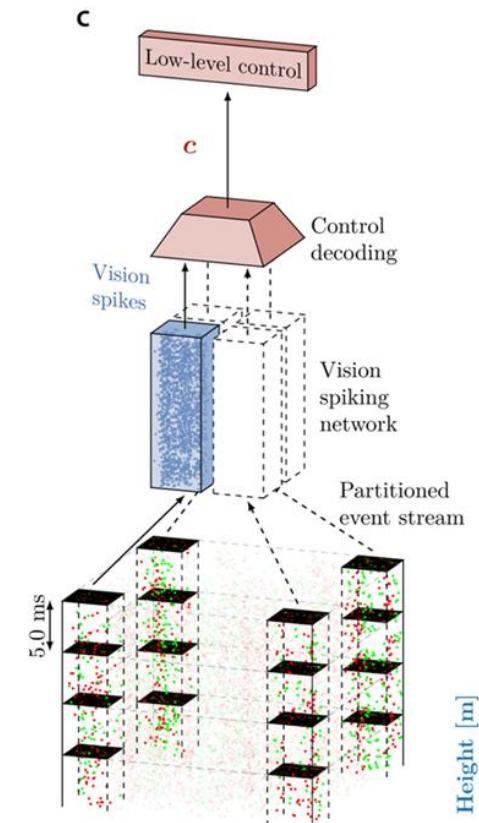
F Paredes-Vall'es, J J Hagenaars, J Dupeyroux, S Stroobants, Y Xu, and G CHE de Croon,
"fully neuromorphic vision and control for autonomous drone flight",
Science Robotics, Vol. 9, No. 90, eadi0591, (2024)

■ Key point

- This paper propose an **SNN model** that outputs **low-level control** commands from **event camera** inputs.

■ Network Architecture

- Divided into two blocks
 1. Vision spiking network (**SNN**)
 - Input : [Event stream (**Event camera**)]
 - Output : [Ego-motion (translational velocity, angular velocity)]
 2. Control decoder (Linear map)
 - Input : [Ego-motion, Target Ego-motion]
 - Output : [Thrust, Torque]



3. Neuromorphic Control using Event-based Vision

11

■ What's Event Camera

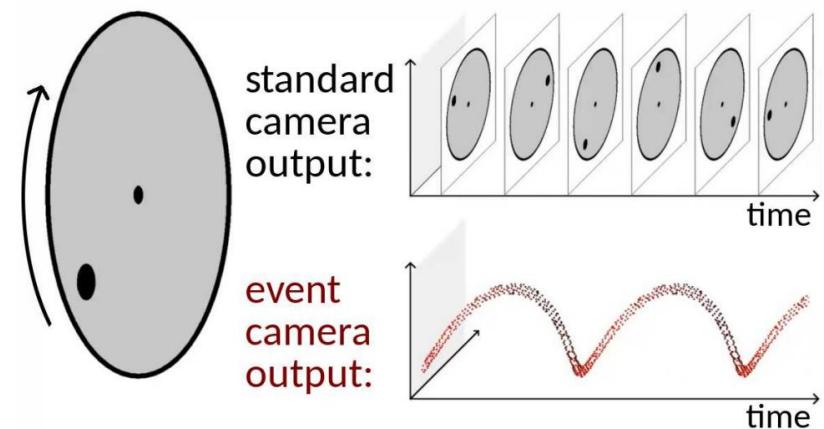
- **Asynchronously** detects changes in brightness for each pixel and outputs them as events
↔ Records brightness values for all pixels **simultaneously** at each time step (RGB camera)

■ Features of Event Camera

- Temporal resolution
- Low power consumption
- High Dynamic Range (HDR)

■ Advantages of combining with SNN

- **Low latency response** through asynchronous processing
- Efficient processing of sparse event data, enabling **low power consumption**
- Natural handling of temporal information



Event camera operation [1]

[1] Gehrig+ "Asynchronous, Photometric Feature Tracking using Events and Frames" @ECCV 2018

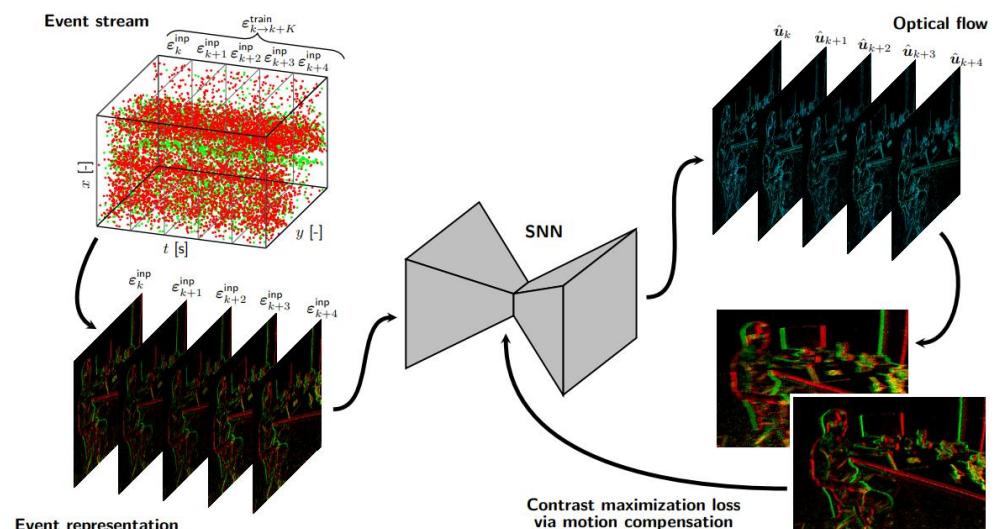
3. Neuromorphic Control using Event-based Vision

■ Optical-flow

- Vectors representing the movement (or motion) of objects and textures between consecutive images

■ Event-based Optical flow

- Implemented as a self-supervised learning
- Using an encoder-decoder model
- Accuracy comparable to ANNs
The model size is too large to be directly implemented on a neuromorphic chip
- Techniques to reduce the network size are required



J Hagenaars, F Paredes-Vall'es, and G De Croon, "self-supervised learning of event-based optical flow with spiking neural networks", Advances in Neural Information Processing Systems, Vol. 34, pp. 7167–7179, 2021.

3. Neuromorphic Control using Event-based Vision

13

■ Vision Network Architecture

— Input (Event camera) / Event buffer

- Uses corner regions instead of the entire camera image
- Temporal accumulation of events to create an event grid

— Encoder Network (SNN)

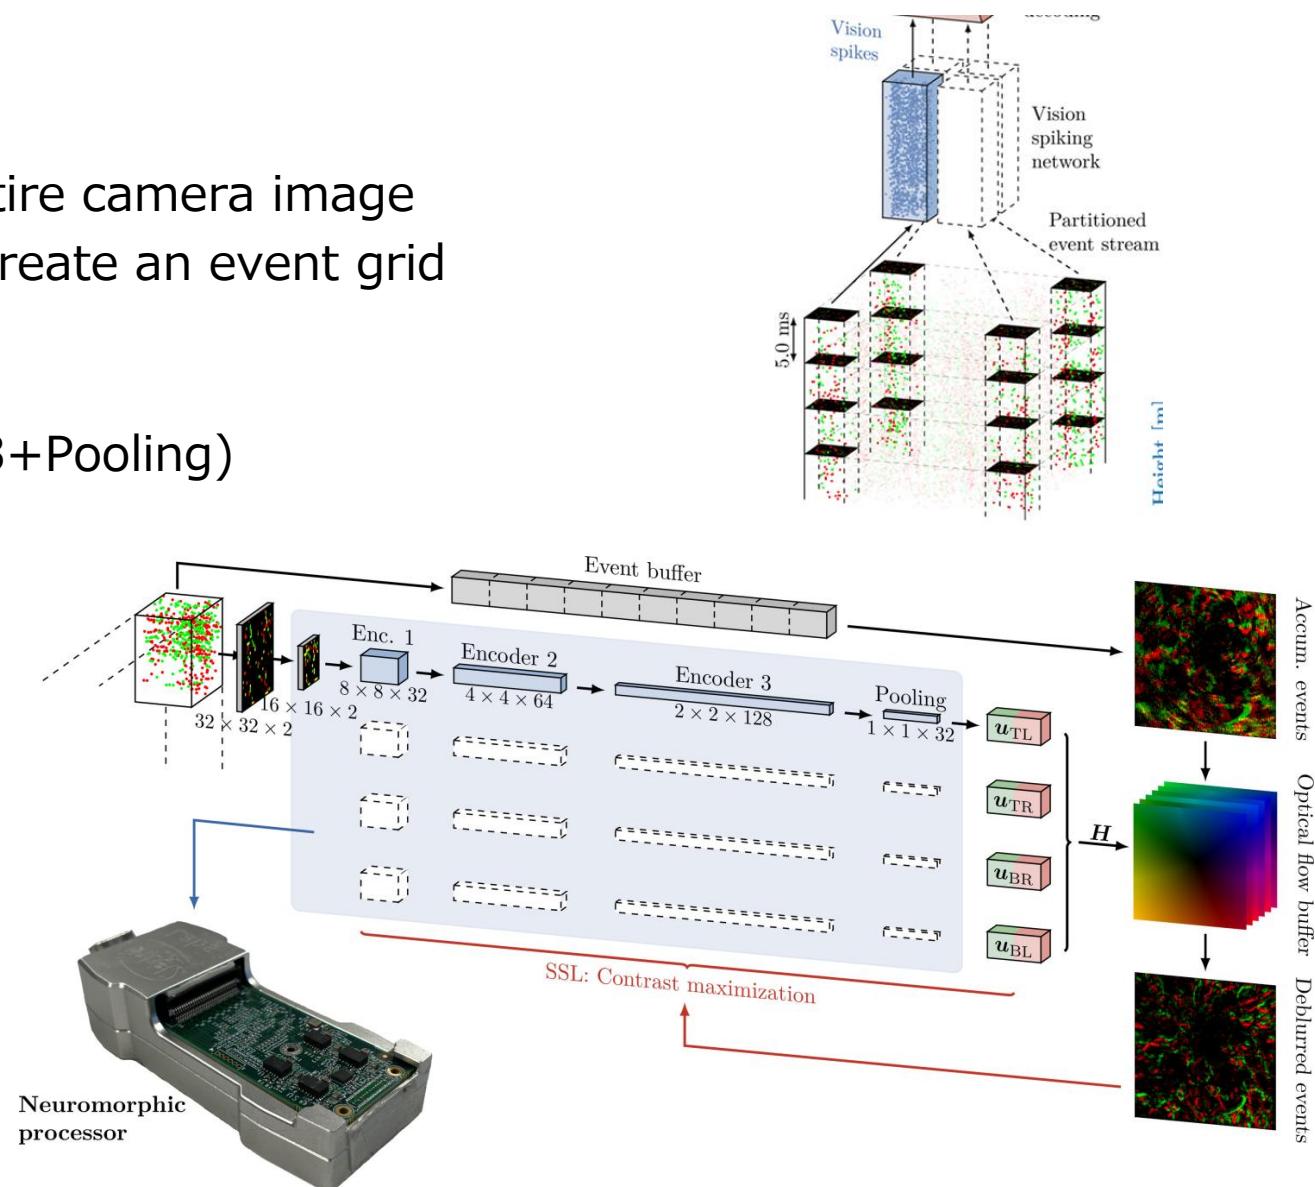
- Output the feature vector (Encoder \times 3+Pooling)

— Optical flow estimation

- Estimating local optical flow from feature vectors
- Combine four local optical flows to create optical flow buffer.

— Output

- Calculating egomotion from optical flow buffer



3. Neuromorphic Control using Event-based Vision

Controller (Linear map, 9×4)

[3DF velocity, 3DoF acceleration, Pitch, Roll, Target velocity]

[Thrust, Roll, Pitch, Yaw]

Evolutionary algorithm

100 random 9×4 matrix

×25000

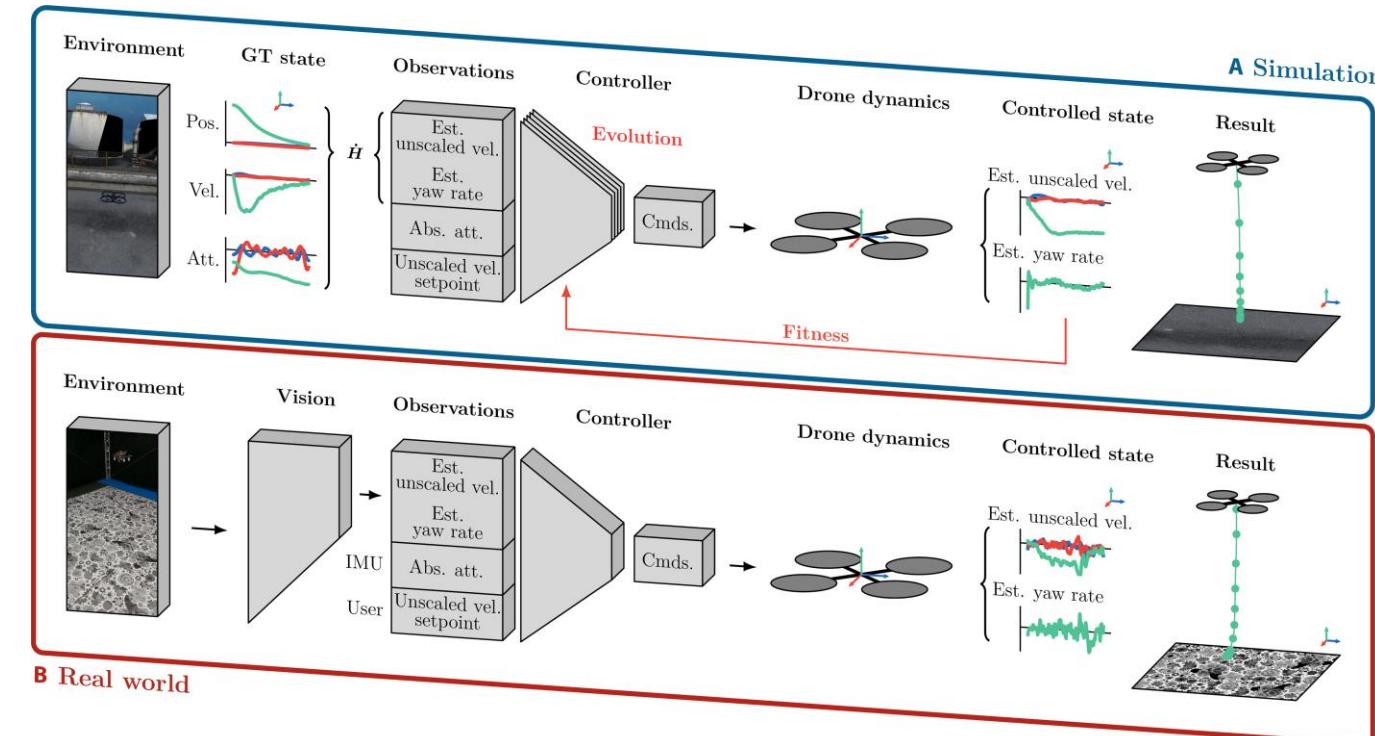
Select the 100 items
with the highest fitness

Fitness function

$$F = \frac{1}{N_{\text{eval}}} \sum_{i \in N_{\text{eval}}} \sum_{j \in N_{\text{steps}}} \mathbf{w} \cdot \left(v_{\text{sp}, i}^{\mathcal{B}} - \begin{bmatrix} \hat{V}_x^{\mathcal{B}} \\ \hat{V}_y^{\mathcal{B}} \\ \hat{V}_z^{\mathcal{B}} \end{bmatrix}_j \right)^2 + (\hat{\omega}_z^{\mathcal{B}})^2$$

Translational velocity error

Yaw penalty



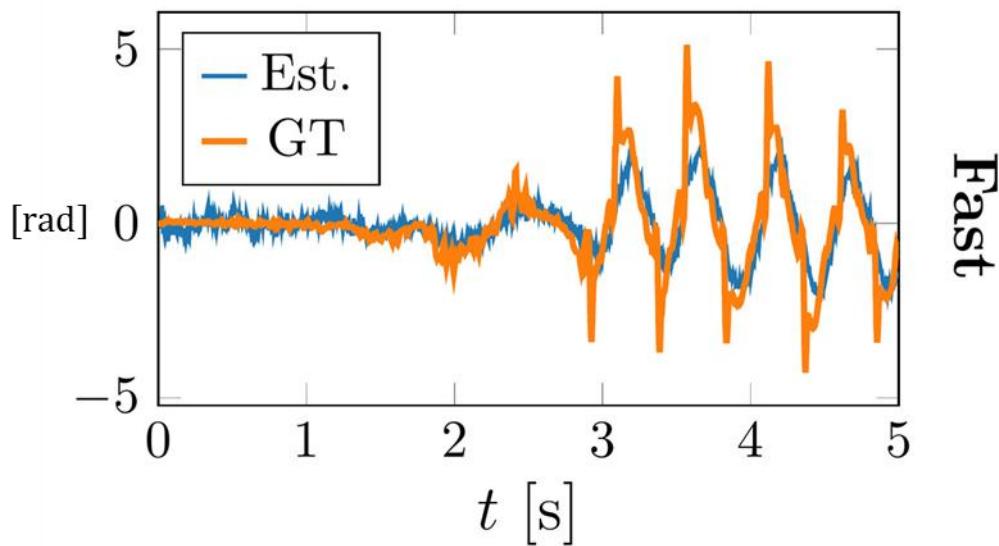
Result : Vision network

- **Condition**

- Training data : Event sequence in a real experimental environment
 - GT : Records of the drone's position and attitude during event data collection

- **Yaw rotation estimation**

- It was able to accurately follow even at a high rotation speed of 4 rad/s.



3. Neuromorphic Control using Event-based Vision

16

■ Result : Evaluation of control

■ Simulation

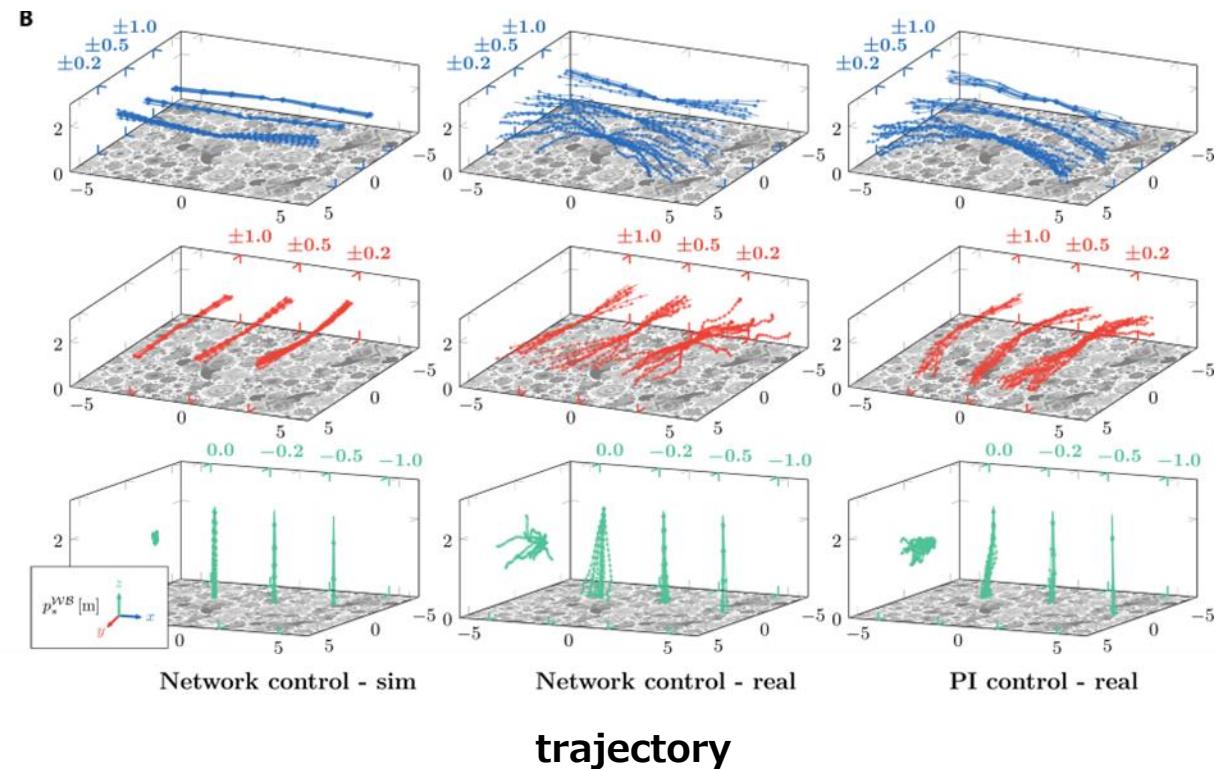
- Success in all flight patterns

■ Real

- Failed in many patterns
 - Reality gap
 - Insufficient integral control

■ Trajectory stability at high speed

- Reduction in noise ratio by increasing the number of events



3. Neuromorphic Control using Event-based Vision

17

■ Power Consumption

- Comparison of SNN execution devices
 - Loihi : 0.95 W / Jetson : 2.98 W

→ Power saving achieved by neuromorphic chips.

→ Achieving greater power savings by reducing **idle power consumption**.

■ Inference speed

- SNN : 1637 [inf/s] / ANN : 83 [inf/s]

→ SNNs enable significantly **faster inference**.

Device	Seq.	Static (W)	Dynamic (W)	Idle (W)	Running (W)	Delta (W)	inf/s	mJ/inf
Nahuku 32 (empty)	Any	0.86	0.04	0.90	0.90	4 × 10 ⁻³	60,496	71 × 10 ⁻⁶
	Slow	0.90	0.05	0.94	0.95	12 × 10 ⁻³	1,637	7 × 10 ⁻³
	Medium	0.90	0.04	0.94	0.95	8 × 10 ⁻³	411	21 × 10 ⁻³
	Fast	0.90	0.04	0.94	0.95	7 × 10 ⁻³	274	27 × 10 ⁻³
3*Nahuku 32: SNN	Slow	–	–	1.04	2.98	1.93	26	75.25
	Medium	–	–	1.06	2.98	1.92	26	75.35
	Fast	–	–	1.04	2.99	1.95	25	76.52
3*Jetson Nano: SNN (10 W)	Slow	–	–	1.04	3.30	2.27	83	27.36
	Medium	–	–	1.07	3.33	2.26	80	28.09
	Fast	–	–	1.07	3.30	2.23	80	27.80
3*Jetson Nano: ANN (10 W)	Slow	–	–	1.04	3.30	2.27	83	27.36
	Medium	–	–	1.07	3.33	2.26	80	28.09
	Fast	–	–	1.07	3.30	2.23	80	27.80

4. Agile Navigation using Reinforcement Learning-Trained SNN

18

Paper3

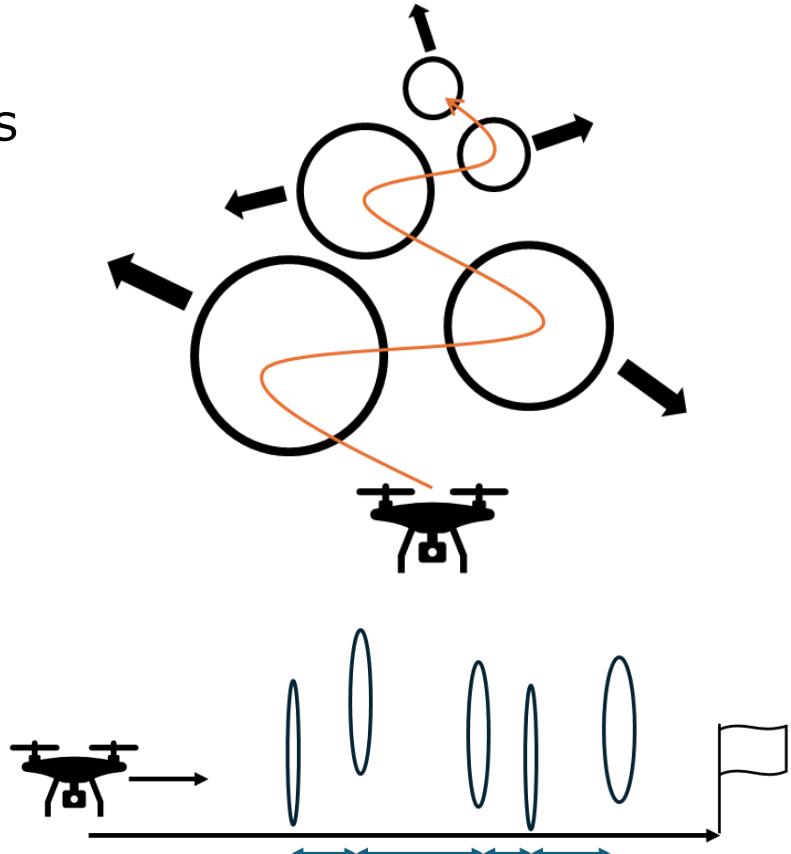
Y-C. Lee, S. Mengozzi, L. Zanatta, A. Bartolini, A. Acquaviva, and F. Barchi,
"Bio-inspired drone control: A reinforcement learning trained spiking neural networks for agile
navigation in dynamic environment", In 2025 IEEE International Conference on Omni layer
Intelligent Systems (COINS), pp. 1–8. IEEE, 2025.

■ Key Point

- Control in **dynamic environments** using **SNN** flight policies
trained with **deep reinforcement learning** algorithms

■ Simulation

- Task : Pass through 5 consecutive rapid-transit gates
- Situation :
 - Drone initial speed : $0 \sim 10$ m/s
 - Distance between gates : $0.5 \sim 9$ m
 - Gates angle : $-\pi \sim \pi$
 - Gates oscillates in the YZ plane



■ Network Architecture (SNN agent)

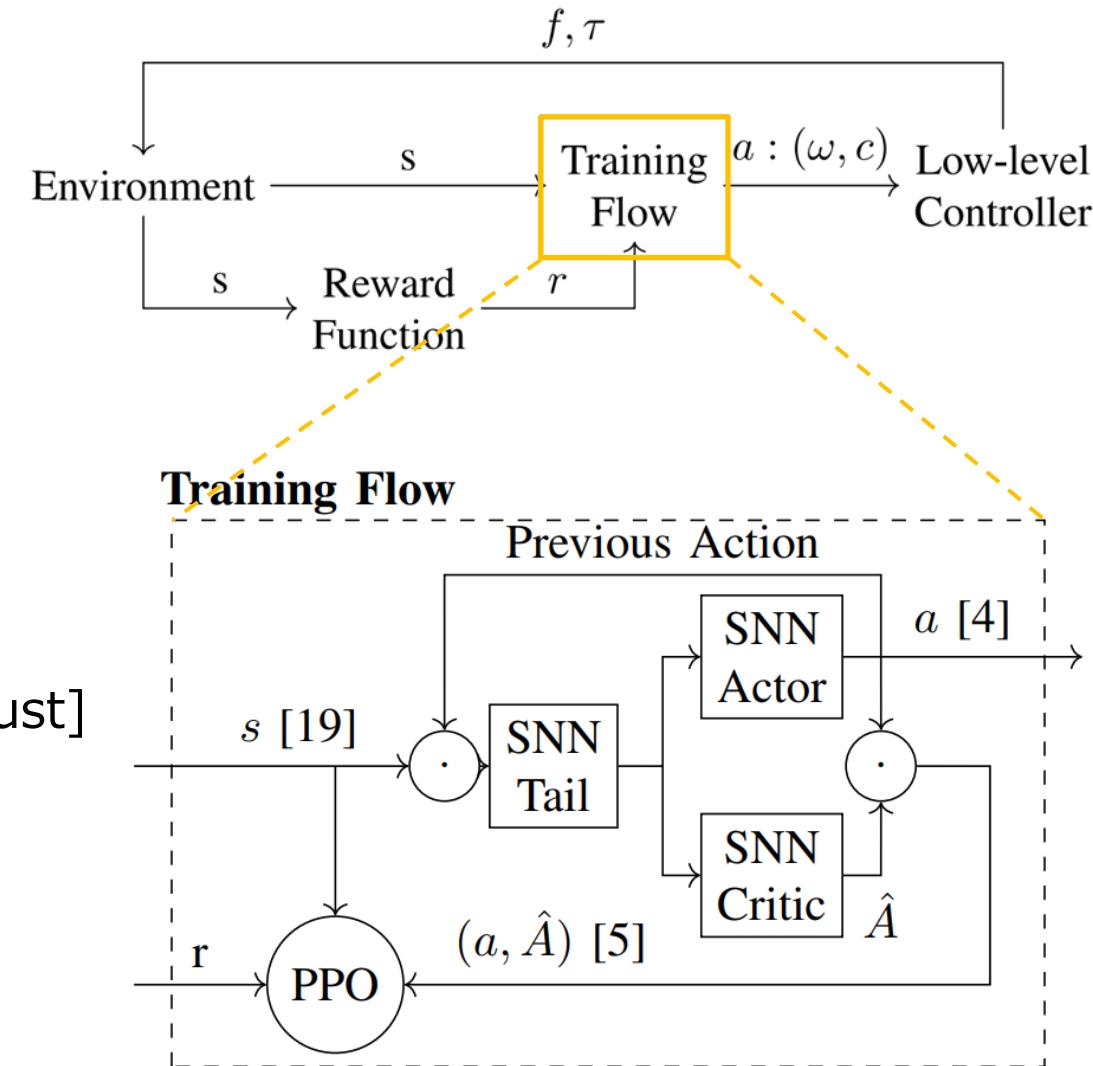
- Input (s):

d_{tgt}	Distance between the drone and the goal
q	Attitude
v	Velocity
ω	Angular velocity
d_{gate}	Distance between the drone and the center of the next gate
v_{gate}	Gates velocity

- Output : [Body Rate (Angular velocity), Thrust]

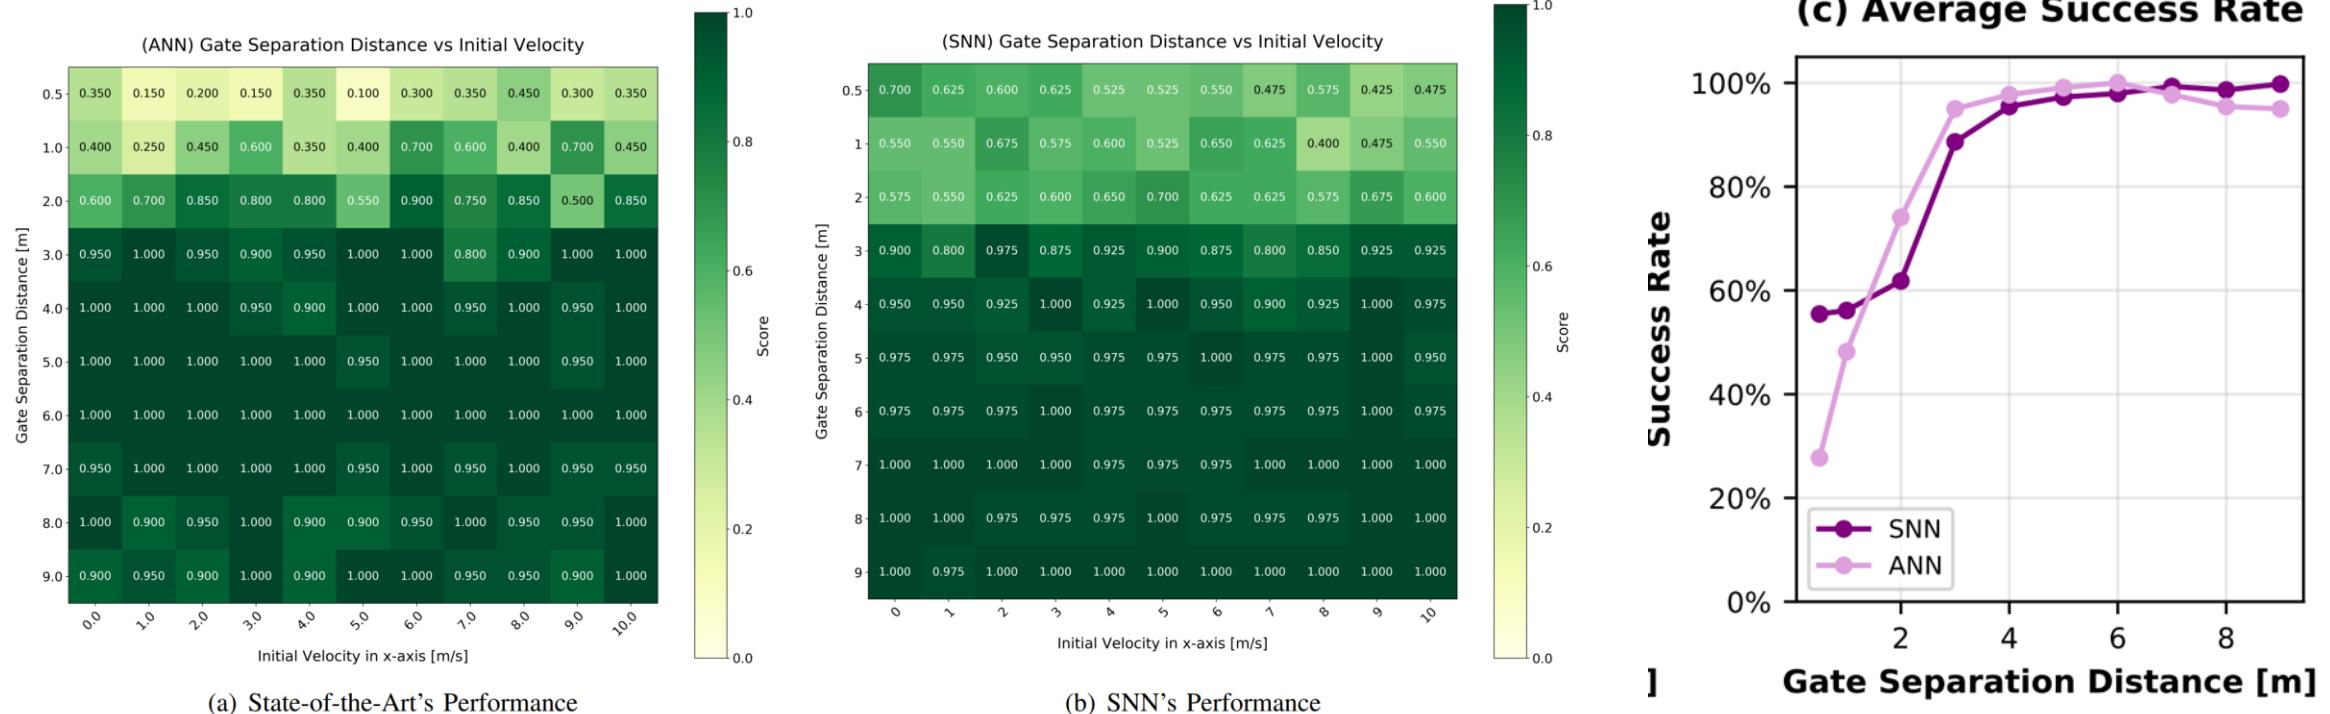
- Model size :

- Input layer : 19 neurons
- Hidden layer : 1024 neurons \times 2
- Output layer : 4 neurons



4. Agile Navigation using Reinforcement Learning-Trained SNN

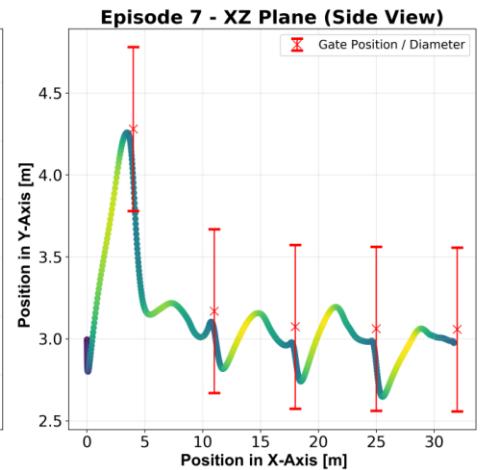
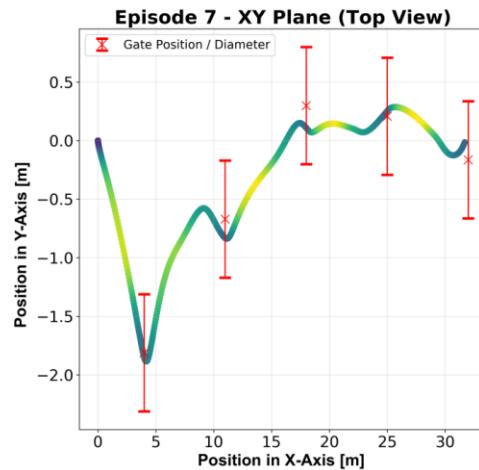
- Result : Gate passing success rate (ANN vs SNN)



- When the gate separation distance is short, SNN is superior
 - This high success rate is due to the SNN's **temporal processing characteristics**

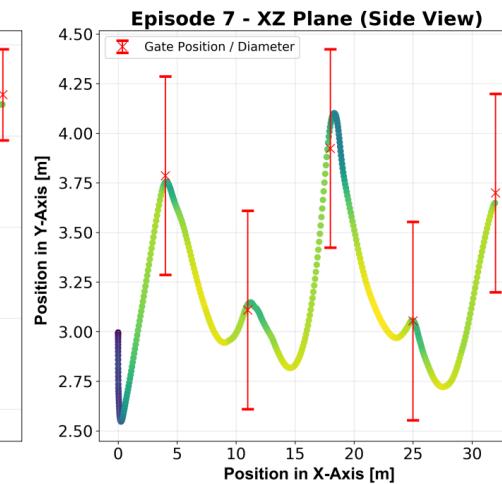
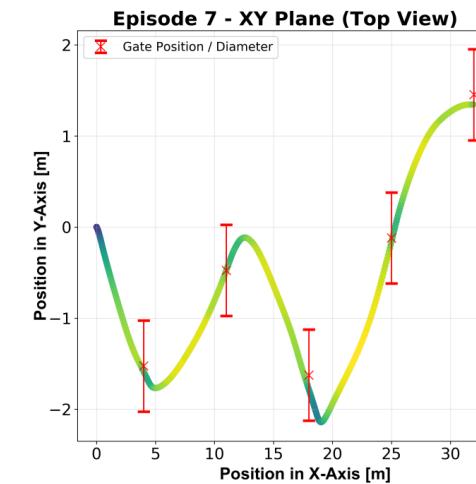
4. Agile Navigation using Reinforcement Learning-Trained SNN

Result : Comparison of trajectories (ANN vs SNN)



Velocity Magnitude [m/s]

(a) State-of-the-Art ANN Trajectory



Velocity Magnitude [m/s]

(b) SNN trajectory

- SNN creates **smoother trajectories**
- SNN consistently maintains **higher speeds** throughout the trajectory
- SNN can complete tasks efficiently in a short amount of time
 - This leads to reduced power consumption

This investigation focused on two applications

low-level control (attitude control) and path planning

■ **low-level control**

- Paper 1 (IMU)
 - Attitude control using SNN with only an IMU is currently less useful than PID control, and integration with image recognition and other technologies is necessary
- Paper 2 (Event camera)
 - There is a problem that the control accuracy deteriorates in response to the demand for reducing the model size.

■ **Path-planning**

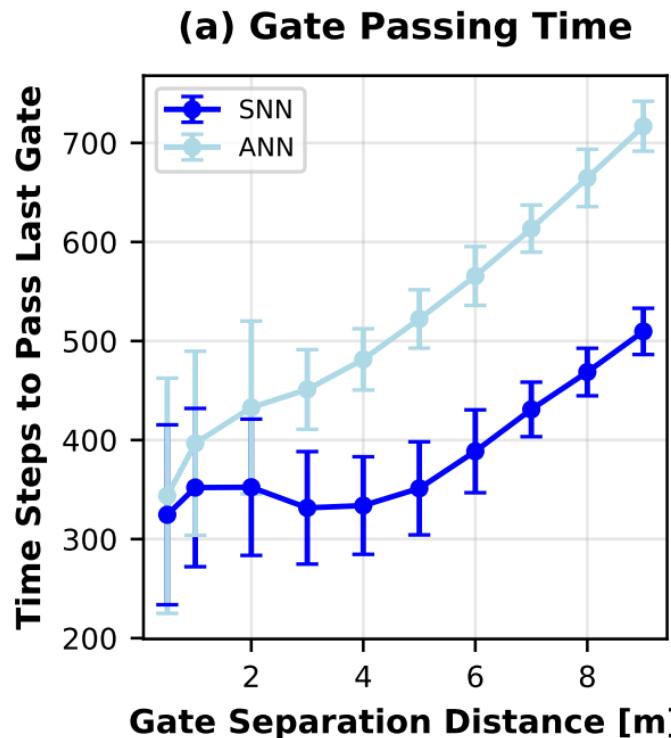
- The challenge remains in designing an optimal reward function for real environments.

■ **Towards Mars exploration**

- While drone control using SNNs has the advantage of reducing power consumption, there are many issues that need to be resolved when applying it to real environments such as Mars.

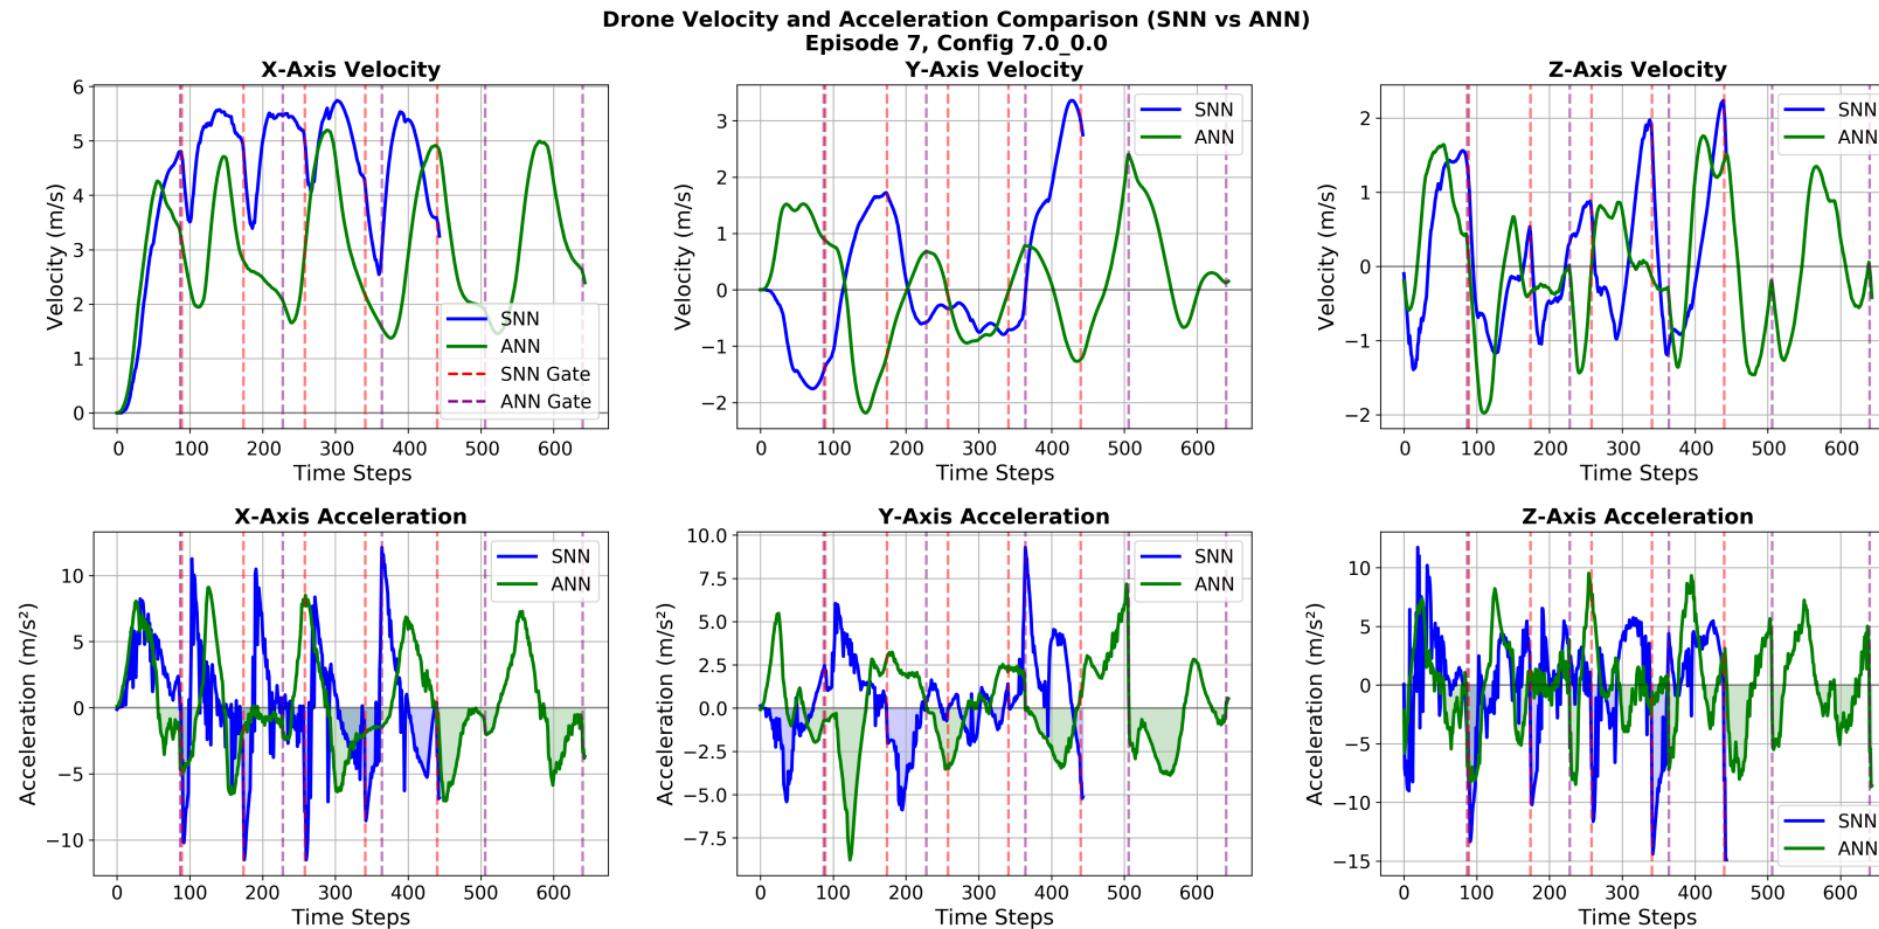
Appendix

- Result : Completion time



4. Agile Navigation using Reinforcement Learning-Trained SNN

25



Paper1

S. Stroobants, C. De. Wagter, and G. CHE. De. Croon,
"Neuromorphic attitude estimation and control",
IEEE Robotics and Automation Letters, (2025).

■ Key point

- End-to-end SNN for attitude estimation and control for real-world micro drone

■ Solution

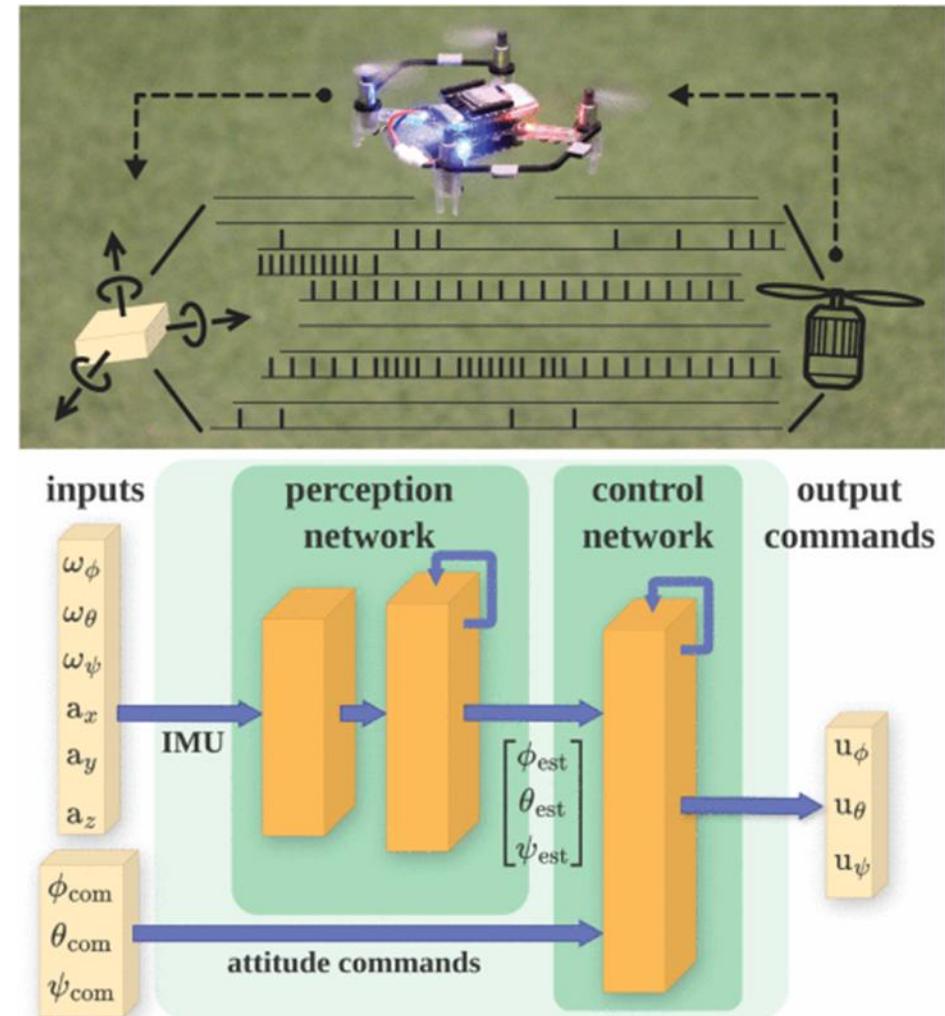
- Division into two subnetworks
- Introduced practical techniques
 1. Time-shifted Targets
 2. Integrator Neurons
 3. Data Augmentation

2. Neuromorphic Attitude Estimation and Control using IMU

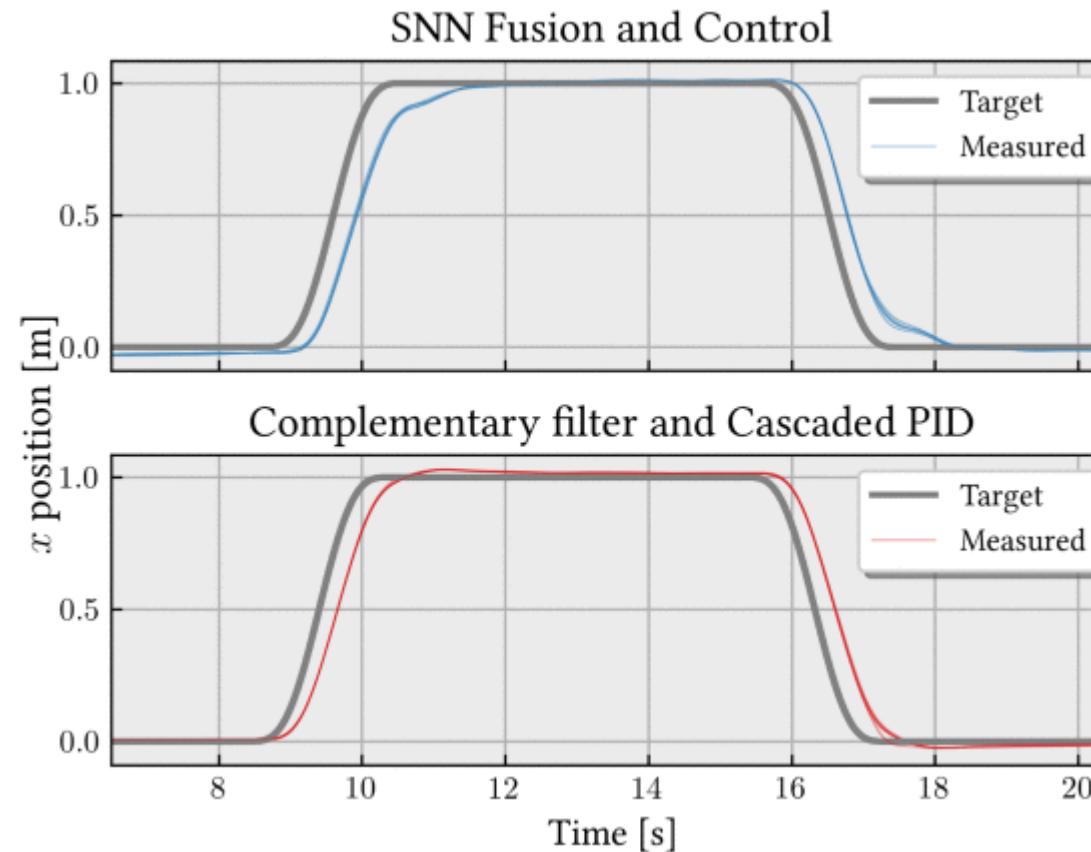
Proposed Network Specifications

Item	Details
Neuron model	CUBA-LIF
Architecture	Perception : 2-layer (recurrent) Control : 1-layer (recurrent)
Number of Nerons	Training : 150-150-130 Inference : 150-100-80
Learning Algorithm	BPTT (Backpropagation Through Time)

- Training data collection
 - Collected from 20 minutes of manual flight
 - Attitude : complementary filter
 - Control : cascaded PID controller
- Spike \Leftrightarrow Real Values
 - Transformation using a Weight Matrix



1. Single-axis position control



SNN can achieve stable target tracking with performance comparable to PID control

Power Consumption estimation

- Number of addition operation (1 step)
 - PID : **3000** times / step
 - SNN : **7500** times / step
 - Energy efficient is expected to be in the **same order**
- **Potential for further power consumption reduction**
 1. Integration with vision tasks
 - Highly compatible with the reduction in multiplications enabled by spike sparsity
 2. Reduction in states like hovering
 - Network sparsity increases when control is not applied

3. Neuromorphic Control using Event-based Vision

31

Paper2

F Paredes-Vall 'es, J J Hagenaars, J Dupeyroux, S Stroobants, Y Xu, and G CHE de Croon,
"fully neuromorphic vision and control for autonomous drone flight",
Science Robotics, Vol. 9, No. 90, eadi0591, (2024)

■ Key point

- Vision-based autonomous flight entirely on a neuromorphic chip

■ Solution

- Implemented a full pipeline on the Loihi chip
 1. **Ego-motion estimation** (optical flow) from event camera data
 2. Velocity control (trained by a genetic algorithm)

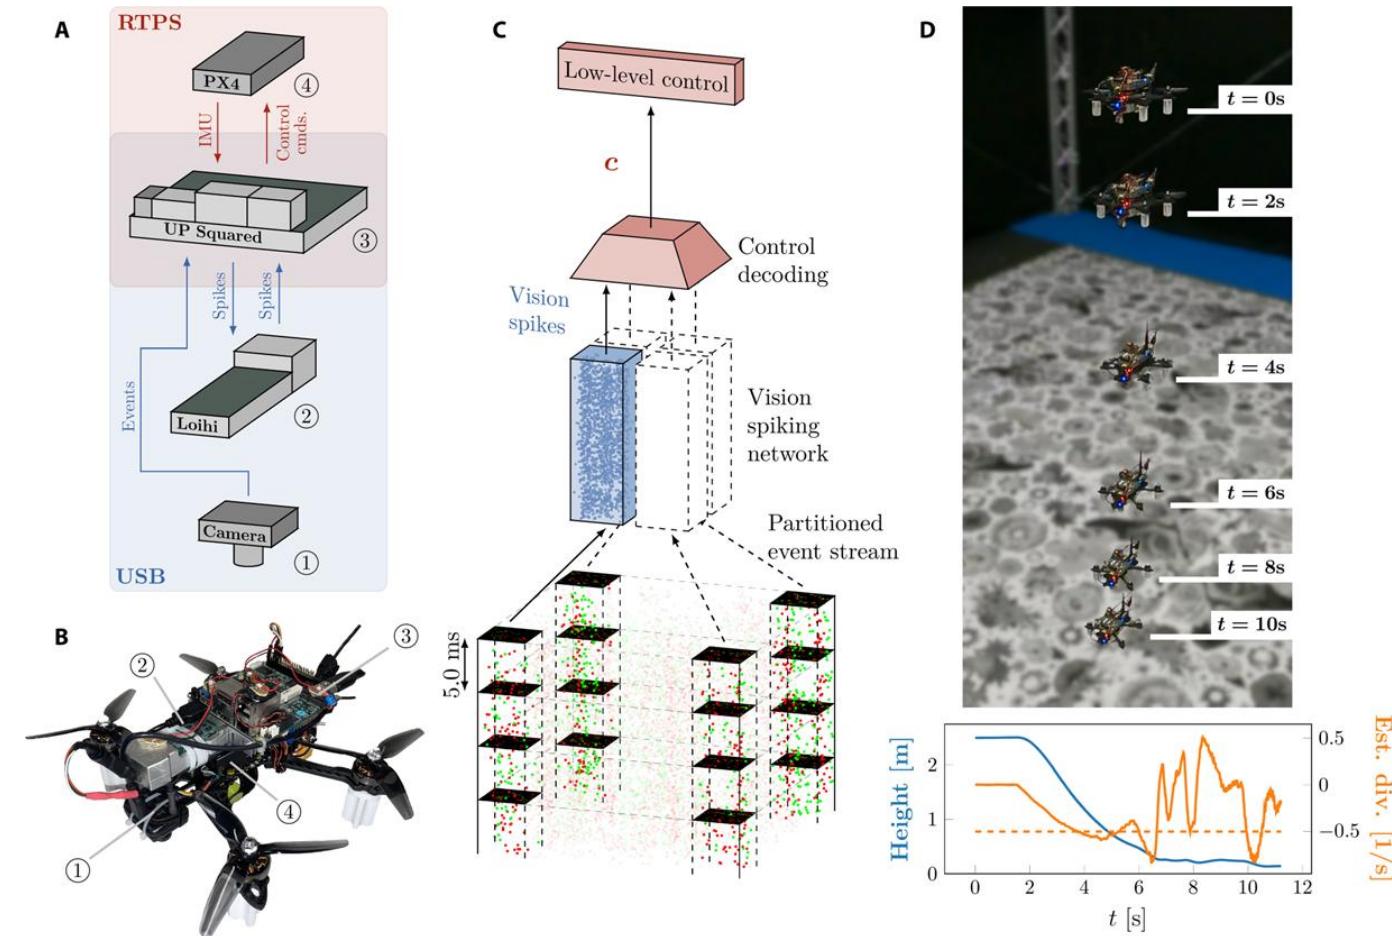
3. Neuromorphic Control using Event-based Vision

- Proposed system architecture

- Vision Network (Optical flow)
- Control decoder (Linear map)

- Network size reduction

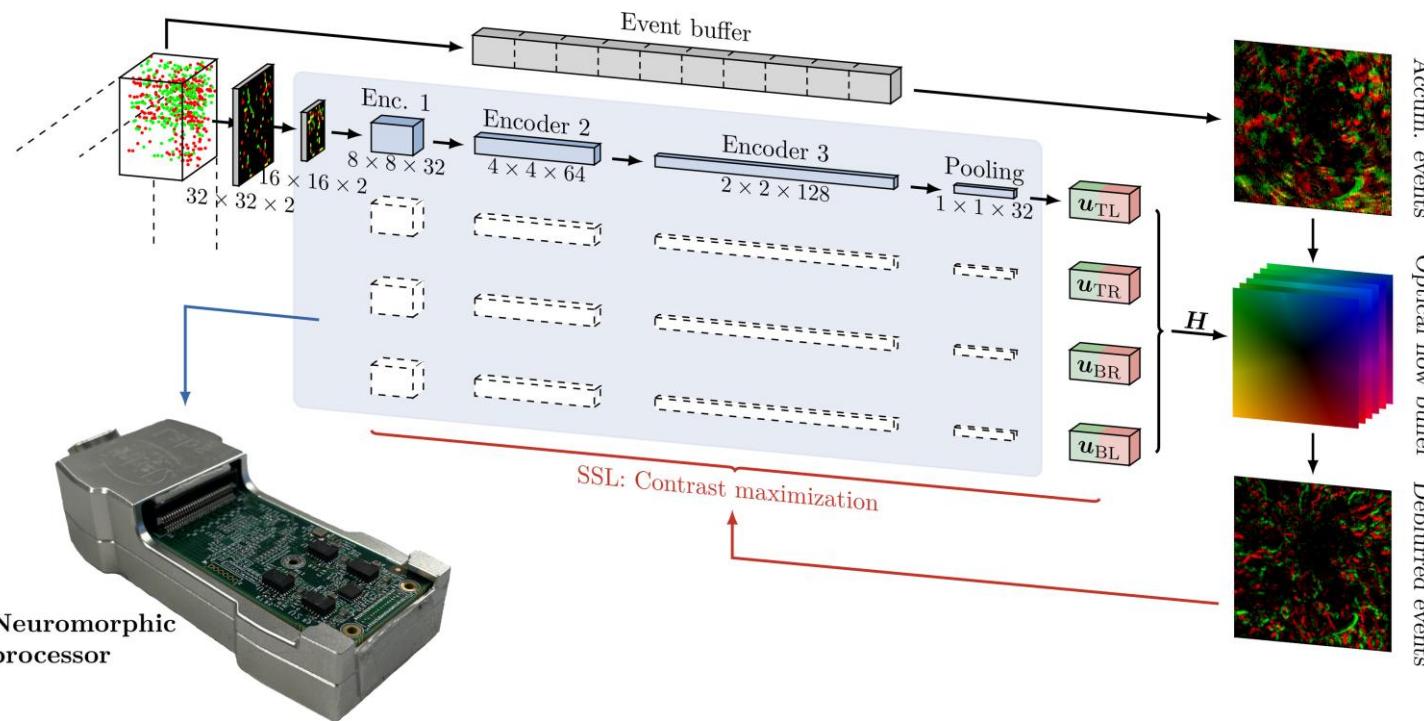
- The floor is a flat surface with a rich texture
- Limit the number of input event (90 per corner ROI)



3. Neuromorphic Control using Event-based Vision

33

Vision Network Architecture



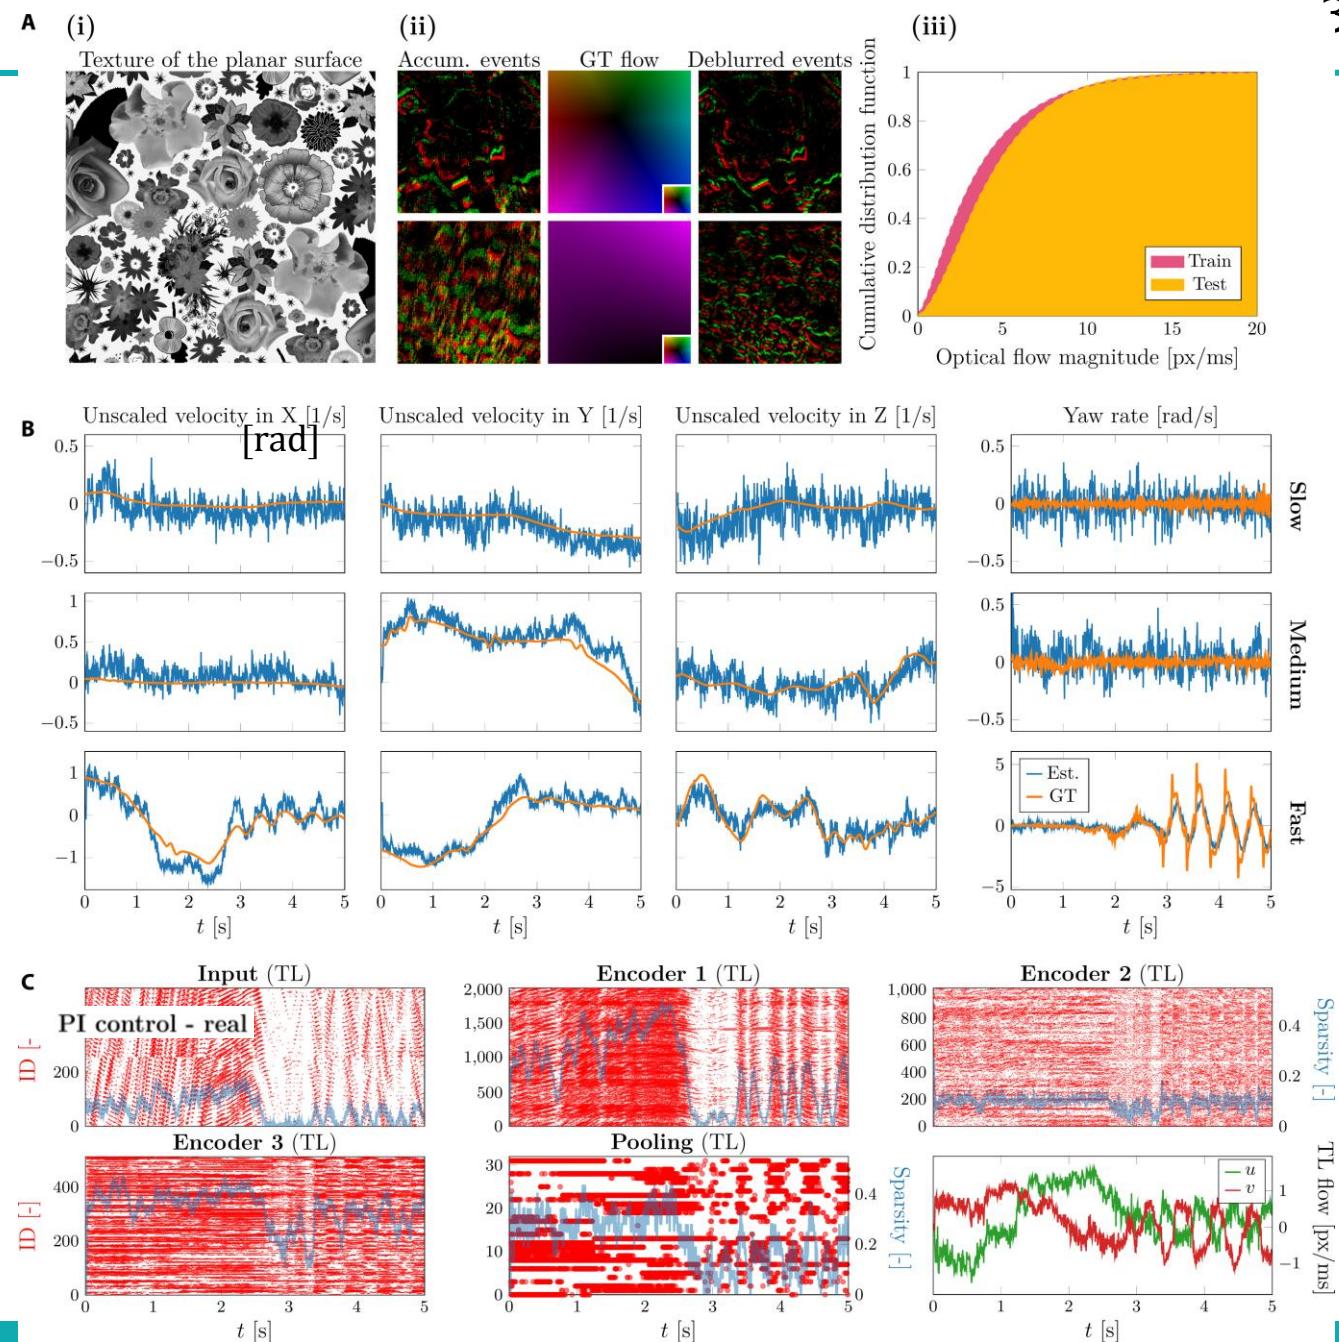
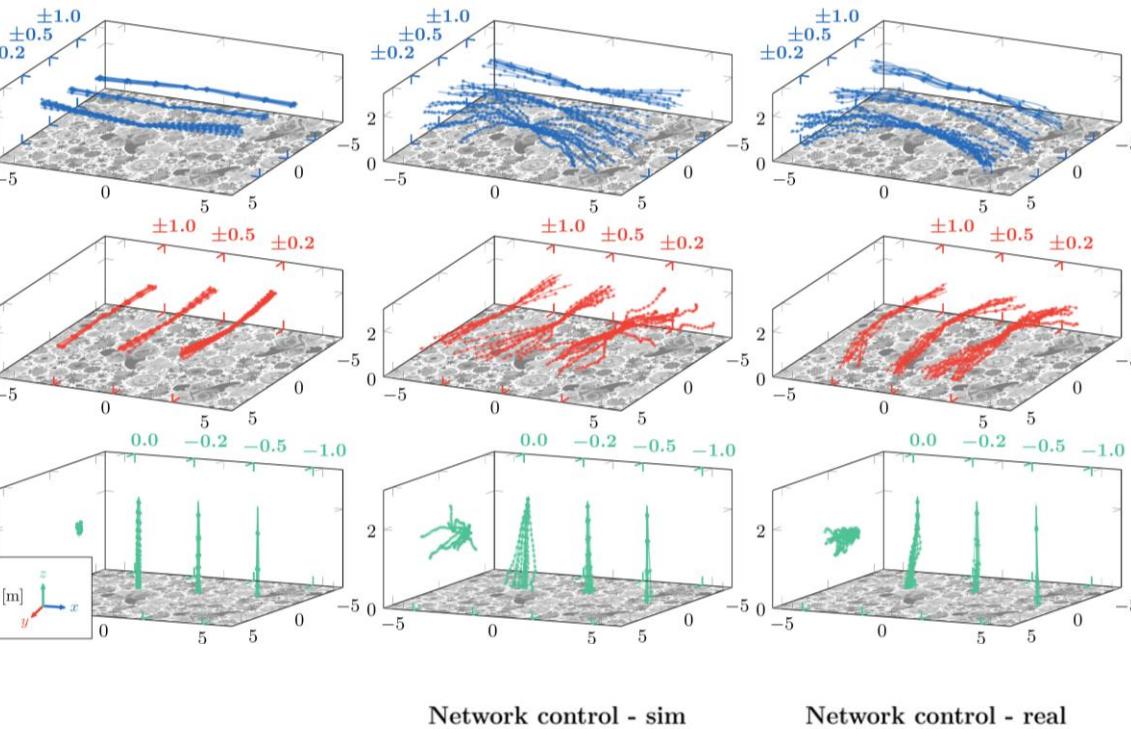
Optical flow estimates

$$\mathbf{u}(\mathbf{x}, \mathbf{H}) = \begin{bmatrix} u(\mathbf{x}, \mathbf{H}) \\ v(\mathbf{x}, \mathbf{H}) \end{bmatrix} = \left(\mathbf{H}, \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} \right)_{\text{Eucl}} - \begin{bmatrix} x \\ y \end{bmatrix}$$

\mathbf{H} : holography matrix

$$L_{\text{flow}} = L_{\text{contrast}} + \lambda L_{\text{smooth}}$$

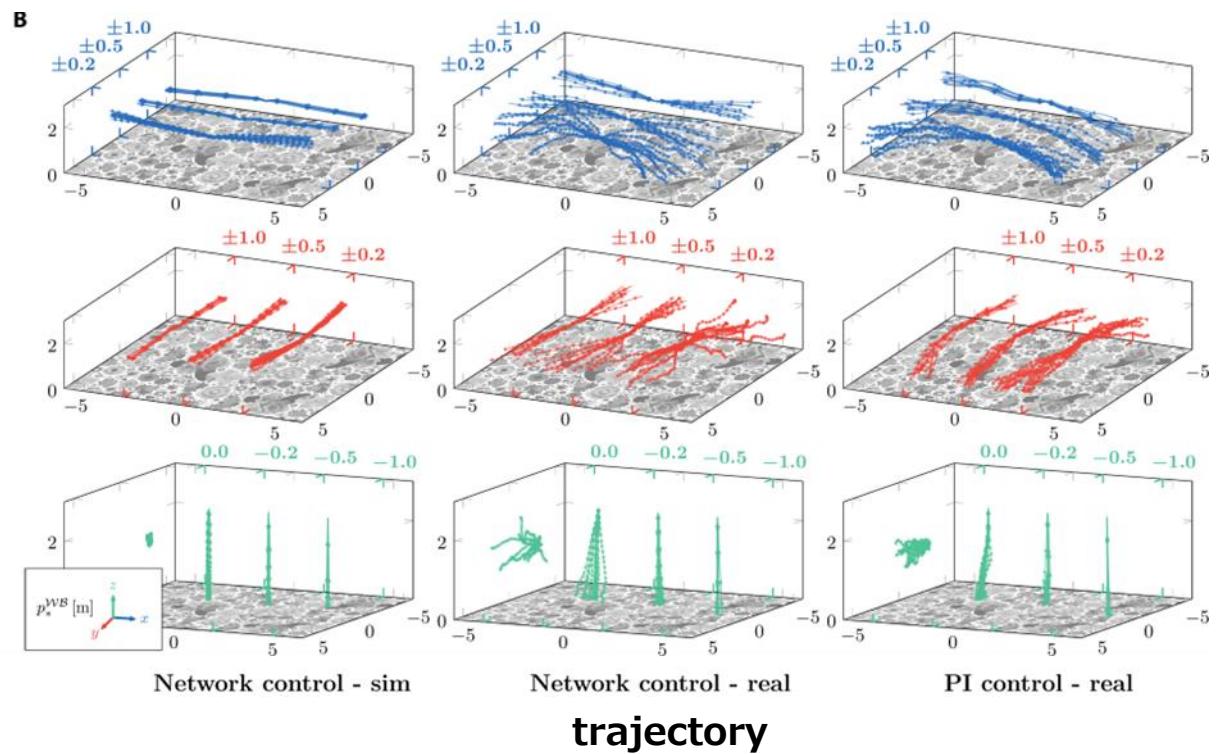
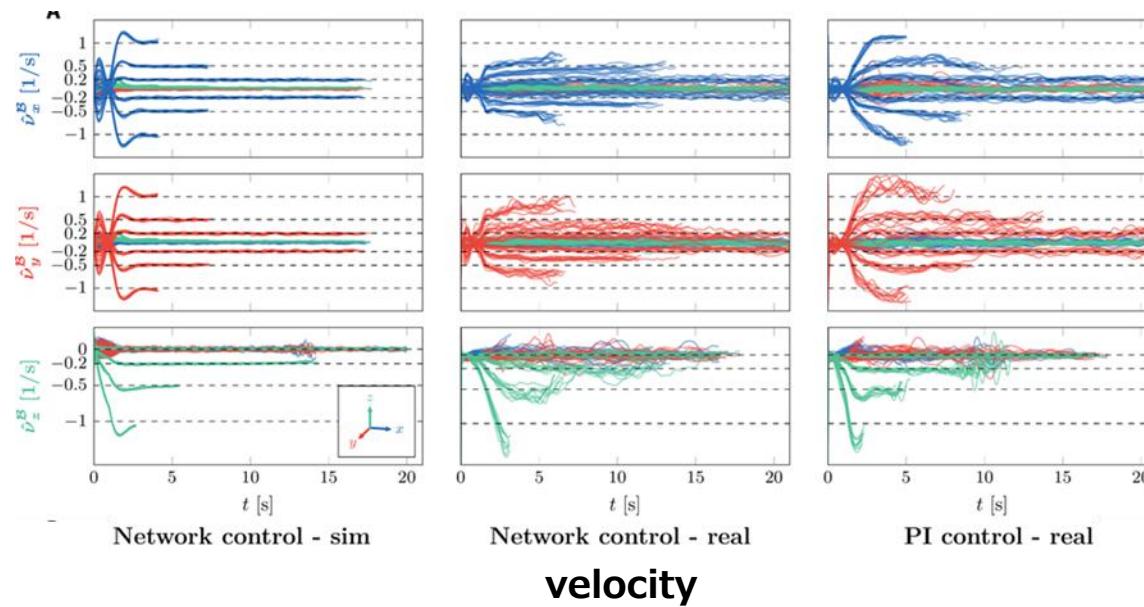
$$\text{EPE} = \frac{1}{N} \sum_{i \in I}^N \|\mathbf{u}_i - \mathbf{u}_i^{\text{GT}}\|$$



3. Neuromorphic Control using Event-based Vision

35

Result : Evaluation of control



- Simulation
 - Success in all flight patterns
- Trajectory stability at high speed
 - Reduction in **noise ratio** by increasing the number of events
- Real
 - Failed in many patterns
Cause → Reality gap / Insufficient integral control

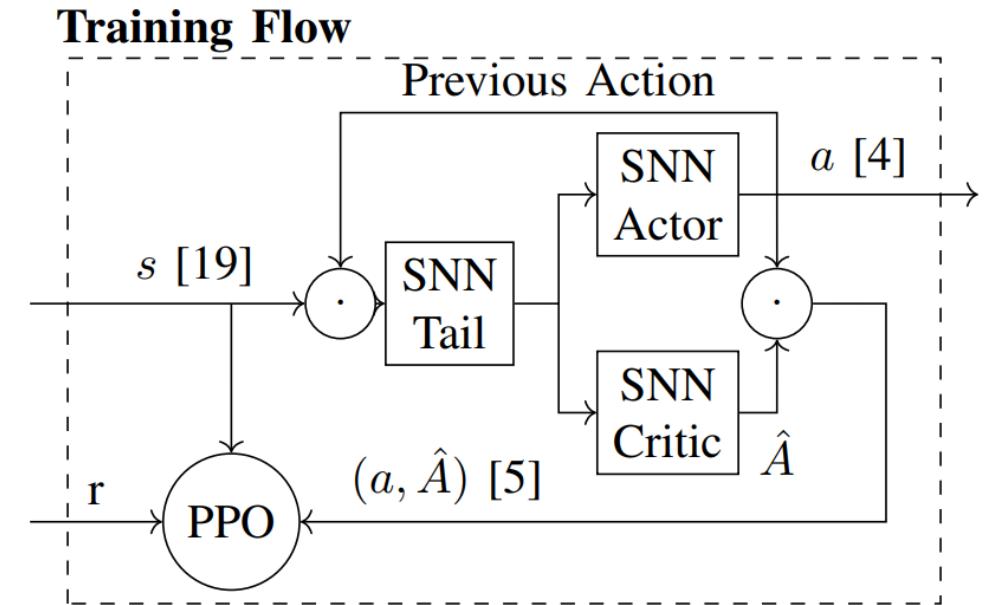
Device	Seq.	Static (W)	Dynamic (W)	Idle (W)	Running (W)	Delta (W)	inf/s	mJ/inf	Device	Seq.	Static (W)	Dynamic (W)	Idle (W)	Running (W)	Delta (W)	inf/s	mJ/inf
Nahuku 32 (empty)	Any	0.86	0.04	0.90	0.90	4×10^{-3}	60,496	71×10^{-6}	3*Jetson Nano: SNN (10 W)	Slow	-	-	1.04	2.98	1.93	26	75.25
3*Nahuku 32: SNN	Slow	0.90	0.05	0.94	0.95	12×10^{-3}	1,637	7×10^{-3}		Medium	-	-	1.06	2.98	1.92	26	75.35
	Medium	0.90	0.04	0.94	0.95	8×10^{-3}	411	21×10^{-3}		Fast	-	-	1.04	2.99	1.95	25	76.52
	Fast	0.90	0.04	0.94	0.95	7×10^{-3}	274	27×10^{-3}	3*Jetson Nano: ANN (5 W)	Slow	-	-	1.05	2.66	1.61	59	27.46
3*Jetson Nano: SNN (5 W)	Slow	-	-	1.05	2.23	1.18	14	86.11		Medium	-	-	1.07	2.64	1.57	56	27.91
	Medium	-	-	1.03	2.25	1.22	14	85.58		Fast	-	-	1.06	2.64	1.58	57	27.52
	Fast	-	-	1.03	2.24	1.21	14	86.19	3*Jetson Nano: ANN (10 W)	Slow	-	-	1.04	3.30	2.27	83	27.36

■ Training Flow

- Input : state of system

$$s \in \mathbb{R}^{19} := (d_{tgt}, q, v, \omega, d_{gate}, v_{gate})$$

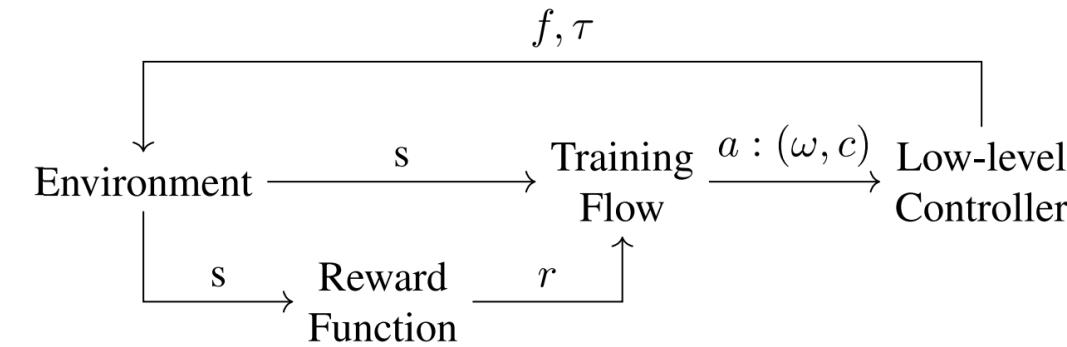
- Output : Thrust, Body rates
- SNN model architecture
 - Hidden layer (1024 neurons) $\times 2$



■ Framework of Reinforcement Learning

- In the proposed framework...
 - Physical simulation environment
 - SNN agent
 - Low-level controller

→ A task of passing through multiple **moving** rings



■ Proximal Policy Optimization(PPO)

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(\underbrace{r_t(\theta) \hat{A}_t}_{\text{Action probability update}}, \underbrace{\text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t}_{\text{Clip function}} \right) \right]$$

- Clipping the policy to **prevent large changes**
→ **Stable learning** is achievable

- Reward function

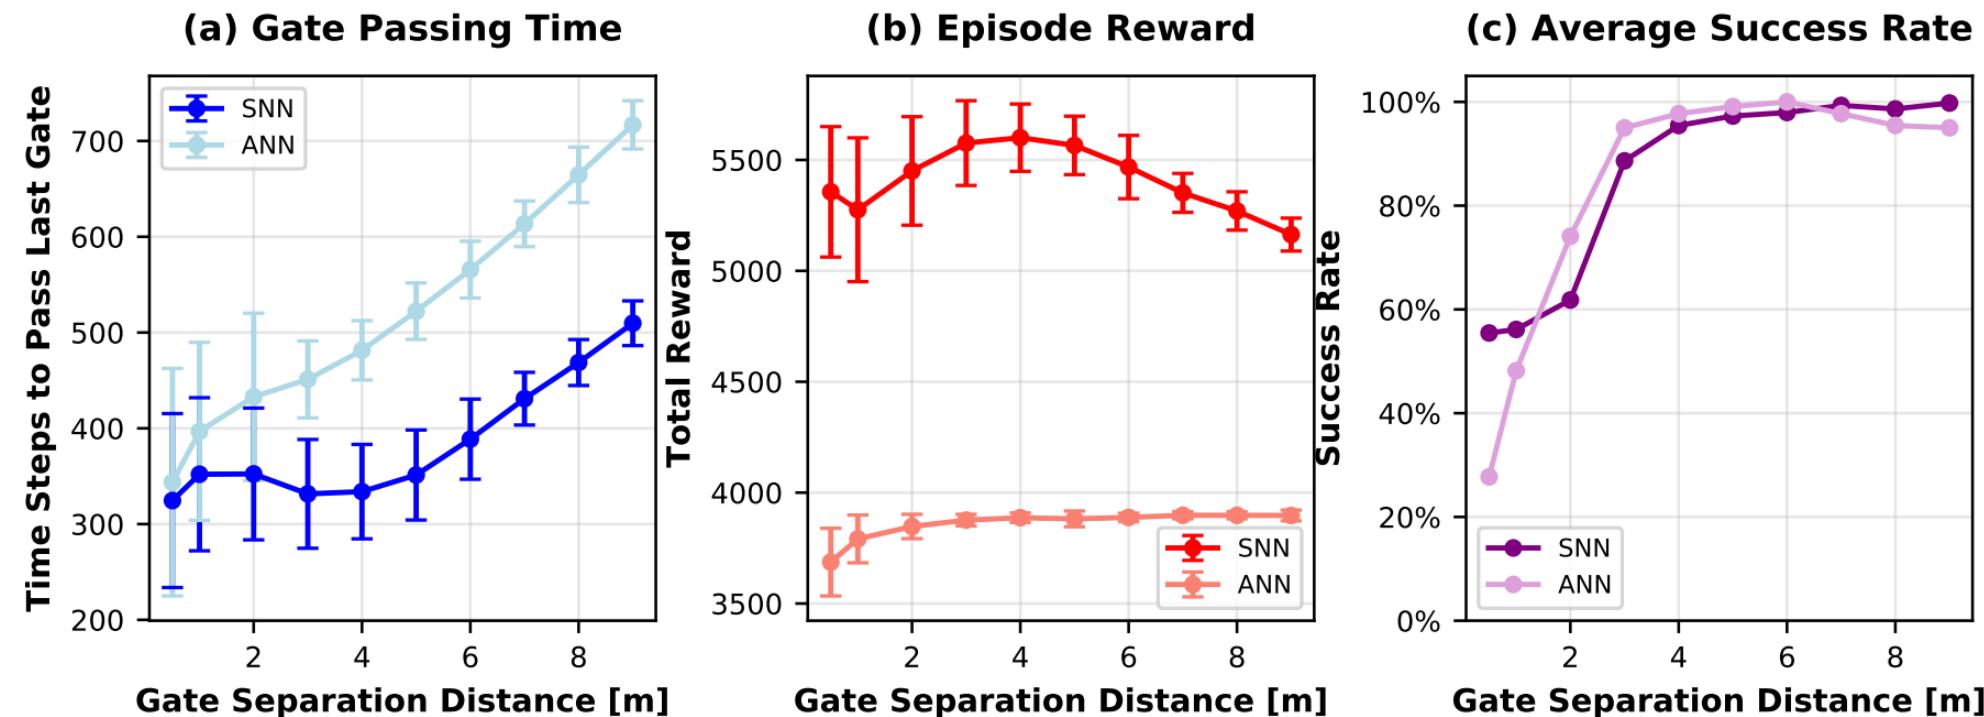
$$r_t = \underline{r_{gate,t}} + \underline{r_{pos,t}} + r_{pos,t} \cdot \underline{(r_{up,t} + r_{spin,t})} - \underline{\beta \cdot r_{acc,t}}$$

Gate passing reward Position reward Horizontal reward Rapid acceleration penalty

$$r_{pos,t} = \frac{0.05}{1 + d_{tgt}^2} + \frac{2.5}{1 + d_{gate}^2} - \frac{2.0}{1 + d_{ring}^2}$$

- Reward
 - 1. Close to final goal point
 - 2. Close to the center of the gate
- Penalty
 - 1. Close to the ring (boundary)

Performance Metrics vs Gate Separation Distance

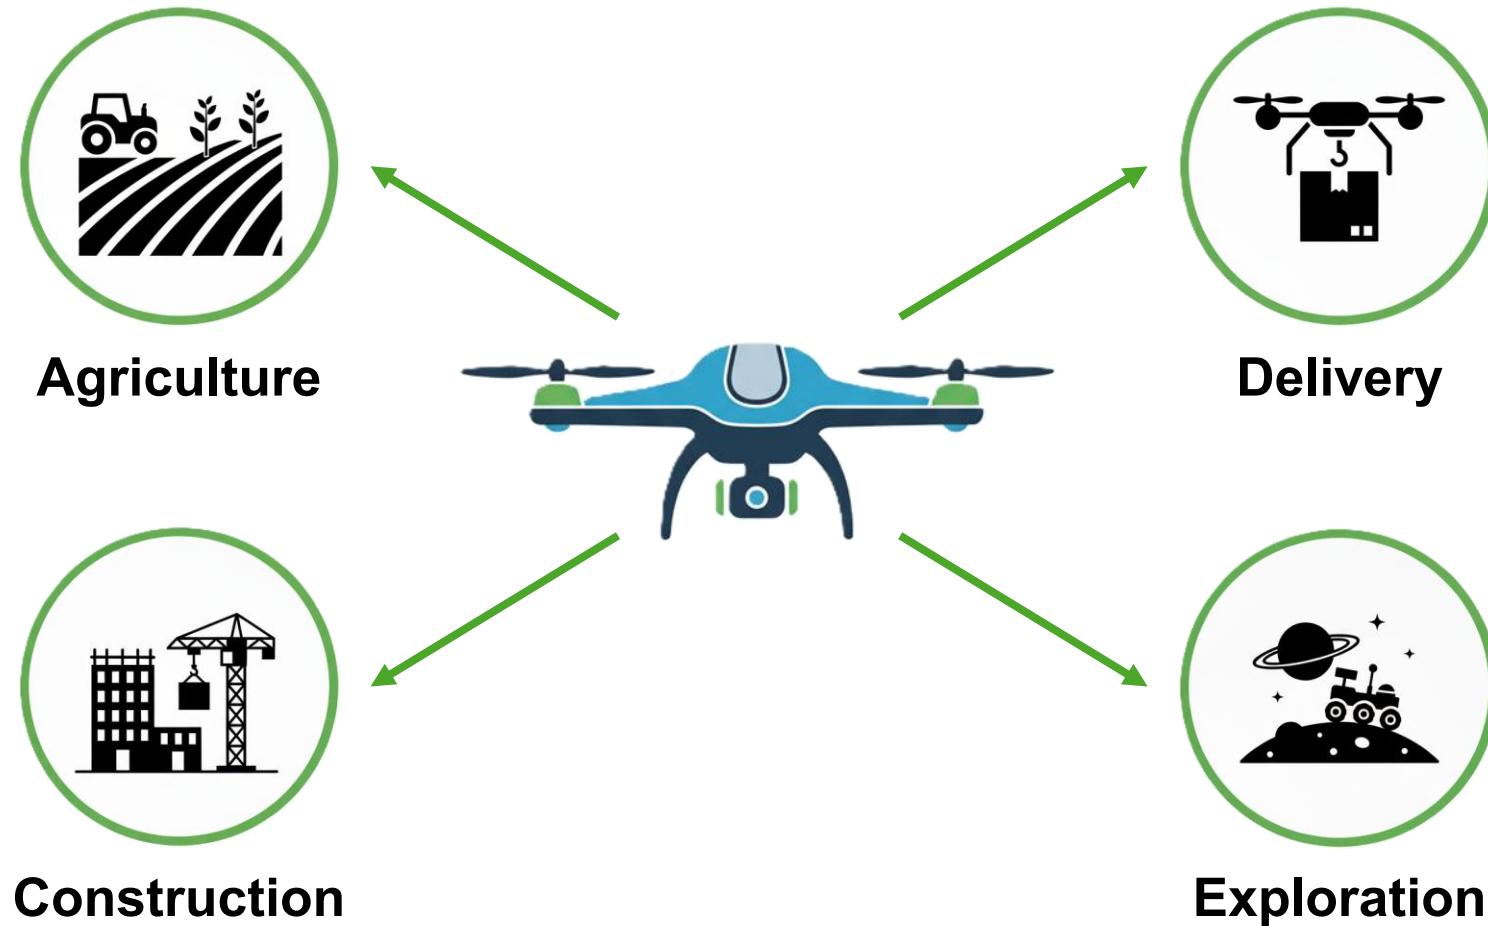


■ Power Consumption

- Number of neurons
 - Input layer : 23
 - Hidden layer : 2048 (1024×2)
 - Output layer : 4
 - Firing rate : 15~20 %
- Number of operation : 215,000

■ Memory usage

- float32 : 4.3 MB
- Int8 : 1.08 MB





On-board

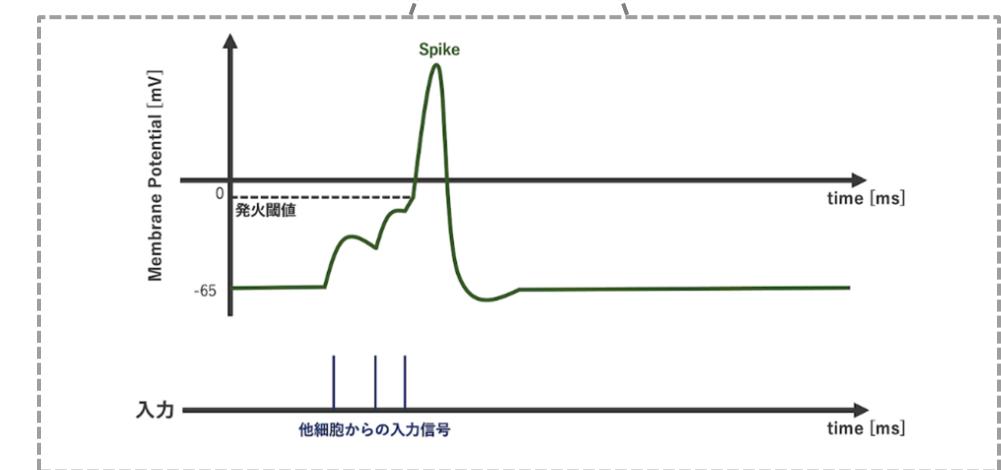
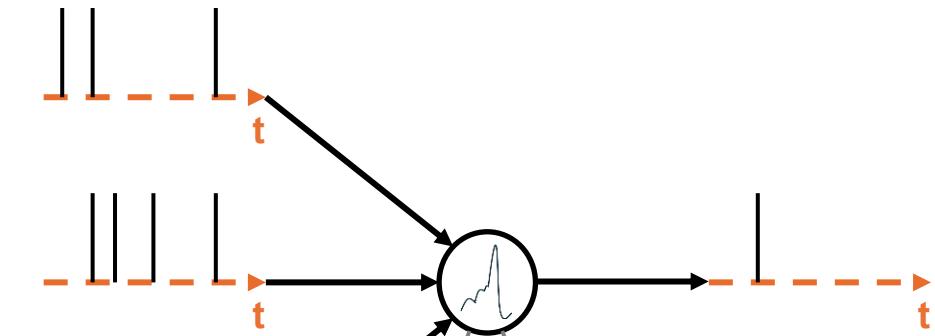
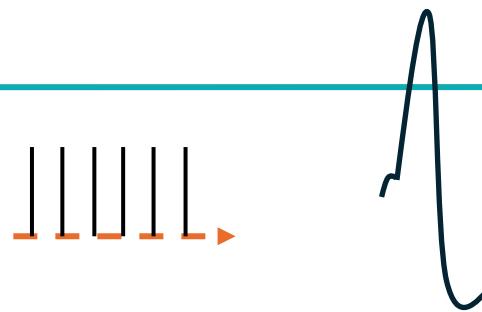
- Real-time processing
- ✗ Power consumption
- ✗ Inference accuracy



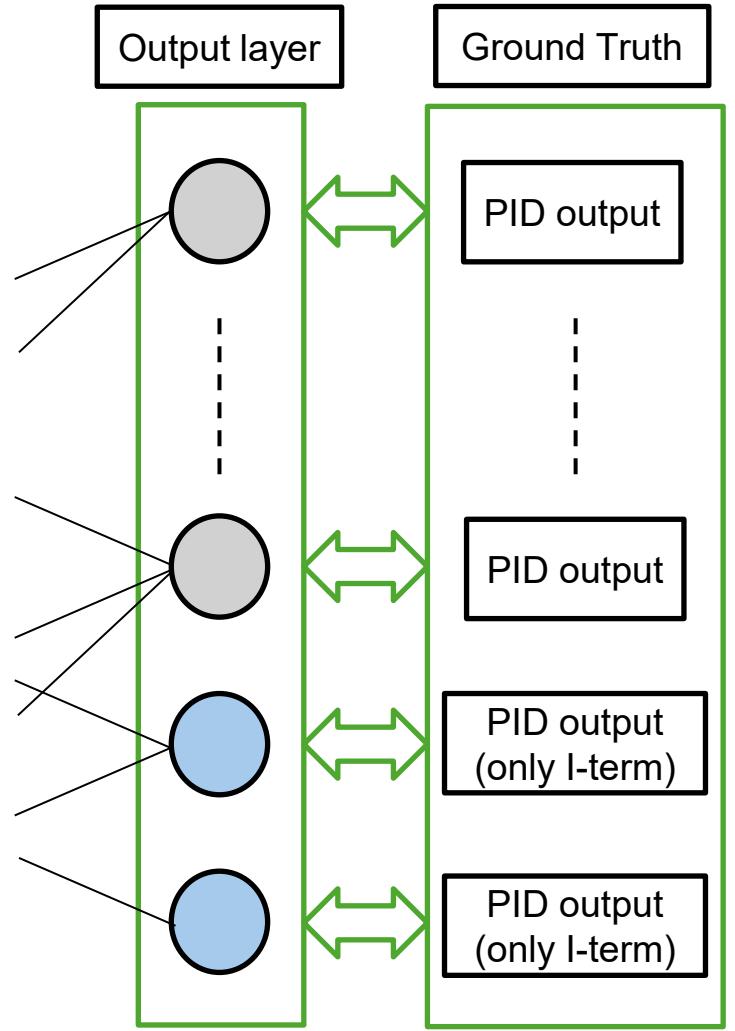
Ground-base

- Inference accuracy
- ✗ Real-time processing









$$RMSE = \sqrt{\frac{1}{n} \sum_{i=0}^{n-1} (y_i - \hat{y}_i)^2}$$

y_i : Ground truth,
 \hat{y}_i : Predicted value,
 n : total data count

